

Contents

1	Introduction	1
1.1	A first look at graphs	1
1.2	Degree and handshaking	6
1.3	Graph Isomorphisms.	9
1.4	Instant Insanity	11
1.5	Trees	15
1.6	Exercises	19
2	Walks	22
2.1	Walks: the basics	22
2.2	Eulerian Walks.	25
2.3	Hamiltonian cycles	30
2.4	Exercises	34
3	Graphs on Surfaces	36
3.1	Introduction to Graphs on Surfaces	36
3.2	The planarity algorithm for Hamiltonian graphs.	40
3.3	Kuratowski's Theorem	45
3.4	Drawing Graphs on Other surfaces	49
3.5	Euler's Theorem	54
3.6	Exercises	58
4	Colourings	61
4.1	Chromatic number	61
4.2	Chromatic index and applications	64
4.3	Introduction to the chromatic polynomial	67
4.4	Chromatic Polynomial continued	71

Chapter 1

Introduction

The first chapter is an introduction, including the formal definition of a graph and many terms we will use throughout. More importantly, however, are examples of these concepts and how you should think about them. As a first nontrivial use of graph theory, we explain how to solve the "Instant Insanity" puzzle.

A first look at graphs

The idea of a graph

First and foremost, you should think of a graph as a certain type of picture, containing dots and lines connecting those dots, like so:




Figure 1.1.1 A graph

We will typically use the letters G , H , or Γ (capital Gamma) to denote a graph. The “dots” or the graph are called *vertices* or *nodes*, and the lines between the dots are called *edges*. Graphs occur frequently in the “real world”, and typically how to show how something is connected, with the vertices representing the things and the edges showing connections.

- *Transit networks:* The London tube map is a graph, with the vertices representing the stations, and an edge between two stations if the tube

goes directly between them. More generally, rail maps in general are graphs, with vertices stations and edges representing line, and road maps as well, with vertices being cities, and edges being roads.

- *Social networks*: The typical example would be Facebook, with the vertices being people, and edge between two people if they are friends on Facebook.
- *Molecules in Chemistry*: In organic chemistry, molecules are made up of different atoms, and are often represented as a graph, with the atoms being vertices, and edges representing covalent bonds between the vertices.




Figure 1.1.2 A Caffeine Molecule, courtesy [Wikimedia Commons](#)

That is all rather informal, though, and to do mathematics we need very precise, formal definitions. We now provide that.

The formal definition of a graph

The formal definition of a graph that we will use is the following:

Definition 1.1.3 A *graph* G consists of a set $V(G)$, called the *vertices* of G , and a set $E(G)$, called the *edges* of G , of the two element subsets of $V(G)$ \diamond

Example 1.1.4 Consider the water molecule, which consists of a single oxygen atom, connected to two hydrogen atoms. It has three vertices, and so $V(G) = \{O, H1, H2\}$, and two edges $E(G) = \{\{O, H1\}, \{O, H2\}\}$ \square

This formal definition has some perhaps unintended consequences about what a graph is. Because we have identified edges with the two things they connect, and have a set of edges, we can't have more than one edge between any two vertices. In many real world examples, this is not the case: for example, on the London Tube, the Circle, District and Picadilly lines all connect Gloucester Road with South Kensington, and so there should be multiple edges between those two vertices on the graph. As another example, in organic chemistry, there are often "double bonds", instead of just one.

Another consequence is that we require each edge to be a two element subset of $V(G)$, and so we do not allow for the possibility of an edge between

a vertex and itself, often called a *loop*.

Graphs without multiple edges or loops are sometimes called *simple graphs*. We will sometimes deal with graphs with multiple edges or loops, and will try to be explicit when we allow this. Our default assumption is that our graphs are simple.

Another consequence of the definition is that edges are symmetric, and work equally well in both directions. This is not always the case: in road systems, there are often one-way streets. If we were to model Twitter or Instagram as a graph, rather than the symmetric notion of friends we would have to work with “following”. To capture these, we have the notion of a *directed graph*, where rather than just lines, we think of the edges as arrows, pointing from one vertex (the source) to another vertex (the target). To model Twitter or Instagram, we would have an edge from vertex a to vertex b if a followed b .

Basic examples and concepts

Several simple graphs that are frequently referenced have specific names.

Definition 1.1.5 The complete graph K_n is the graph on n vertices, with an edge between any two distinct vertices. \diamond

Definition 1.1.6 The empty graph E_n is the graph on n vertices, with no edges. \diamond

Definition 1.1.7 The path graph P_n is the graph on n vertices $\{v_1, \dots, v_n\}$ with edges $\{\{v_1, v_2\}, \{v_2, v_3\}, \dots, \{v_{n-1}, v_n\}\}$. \diamond

Definition 1.1.8 The cycle graph C_n is the graph on n vertices $\{v_1, \dots, v_n\}$ with edges $\{\{v_1, v_2\}, \{v_2, v_3\}, \dots, \{v_{n-1}, v_n\}, \{v_n, v_1\}\}$. \diamond





Figure 1.1.9 Basic graphs

Definition 1.1.10 The complement of a simple graph G , which we will denote G^c , and is sometimes written \overline{G} , is the graph with the same vertex set as G , but $\{v, w\} \in E(G^c)$ if and only if $\{v, w\} \notin E(G)$; that is, there is an edge between v and w in G^c if and only if there is not an edge between v and w in G . \diamond

Example 1.1.11 The empty graph and complete graph are complements of each other; $K_n^c = E_n$.

The path graph P_4 and its complement are shown below:




□

It commonly occurs that there are two different types of vertices, and the edges only go between vertices of the two types. For example, we may be modelling a company, and one type of vertices may represent the employees, and another type of vertices could represent the different jobs that need done, with an edge between a worker and a job if the worker is qualified to do that particular job. We call these graphs *bipartite*.

Definition 1.1.12 A graph \mathbf{G} is *bipartite* if its vertices can be coloured red and blue so that every edge goes between a red vertex and a blue vertex. ◇

Example 1.1.13 The graph below is bipartite.



□

As another example, note that the cycle graph C_4 is bipartite -- we can colour vertices 1 and 3 red, and vertices 2 and 4 blue. But the cycle graph C_3 is *not* bipartite: as the two colours are interchangeable, we can assume we coloured vertex 1 red; then since it is adjacent to both 2 and 3, those vertices must both be blue, but they're adjacent to each other, which violates the definition of bipartite. More generally, we have:

Lemma 1.1.14 *The cycle graph C_n is bipartite if and only if n is even.*

Proof. Let's try to colour the vertices of C_n red and blue so that adjacent vertices have different colour. Without loss of generality, we may assume that v_1 is coloured blue. Then since it is adjacent to v_1 , v_2 must be coloured red. Continuing in this way, we see that v_k is blue for odd k and red for even k . But v_n is adjacent to v_1 , and so these will have different colours precisely when n is even. ■

Lemma 1.1.15 *A graph \mathbf{G} is bipartite if and only if \mathbf{G} has no subgraphs that are isomorphic to C_{2k+1}*

Proof. First, note that if \mathbf{G}_2 is a subgraph of \mathbf{G}_1 , and \mathbf{G}_1 is bipartite, then so is \mathbf{G}_2 : colouring the vertices of \mathbf{G}_1 red and blue in particular colours the vertices of \mathbf{G}_2 as well. Hence, we see if that \mathbf{G} has a subgraph isomorphic to C_{2k+1} , which isn't bipartite by the previous lemma, then \mathbf{G}_1 can't be bipartite, either.

In the other direction, we assume that \mathbf{G} has no subgraphs isomorphic to C_{2k+1} ; we need to prove that \mathbf{G} is bipartite. Again, let's try to colour the vertices of \mathbf{G} red and blue so that adjacent vertices have different colours. Choose a vertex v of \mathbf{G} , without loss of generality we may assume that v is coloured blue; then all vertices adjacent to v -- i.e., those vertices at distance 1 from v -- are coloured red. The vertices adjacent to those must be blue, and the ones adjacent to those must be red, alternating out. This suggests trying to colour all vertices at odd distance from v red, and those vertices at even distance from v blue. We will show that if this colouring has two vertices of the same colour that are adjacent, then \mathbf{G} has an odd cycle.

Assume that u and w are two vertices coloured red that are adjacent. Since u is red, the distance from v to u is odd -- there is a path $v = v_0 - v_1 - \dots - v_{2k+1} = u$. Similarly, there is an odd length path from v to w : $v = w_0 - w_1 - \dots - w_{2\ell+1} = w$. Taking the union of these two subgraphs and the edge connecting u and w , we get a closed walk consisting of $(2k+1) + (2\ell+1) + 1 = 2k + 2\ell + 3$ edges, which is odd. This walk may repeat some vertices and edges, but if it does we can split it into two smaller walks, one of which must have odd length, and eventually we must get a closed walk of odd length that doesn't repeat any vertices.



The case that u and w are both coloured blue is completely analogous, except we will be merging two paths with an even number of edges and one extra edge to get a path with odd length. ■

A special type of bipartite graph is the complete bipartite graphs $K_{m,n}$, which are the simple graphs that have as many edges as possible while still being bipartite.

Definition 1.1.16 The *complete bipartite graph* $K_{m,n}$ is the graph with m red vertices and n blue vertices, and an edge between every red vertex and every blue vertex. ◇

Example 1.1.17 The complete bipartite graph $K_{2,2}$ is isomorphic to C_4 .

The graphs $K_{1,3}$ and $K_{4,4}$ are drawn below.

 $K_{1,3}$  $K_{4,4}$

□

Degree and handshaking

Definition of degree

Intuitively, the *degree* of a vertex is the “number of edges coming out of it”. If we think of a graph G as a picture, then to find the degree of a vertex $v \in V(G)$ we draw a very small circle around v , the number of times the G intersects that circle is the degree of v . Formally, we have:

Definition 1.2.1 Let G be a simple graph, and let $v \in V(G)$ be a vertex of G . Then the *degree of v* , written $d(v)$, is the number of edges $e \in E(G)$ with $v \in e$. Alternatively, $d(v)$ is the number of vertices v is adjacent to. ◇

Example 1.2.2




Figure 1.2.3 The graph K

In the graph K shown in [Figure 1.2.3](#), vertices a and b have degree 2, vertex c has degree 3, and vertex d has degree 1. □

Note that in the definition we require G to be a simple graph. The notion of degree has a few pitfalls to be careful of G has loops or multiple edges. We still want to the degree $d(v)$ to match the intuitive notion of the “number of edges coming out of v ” captured in the drawing with a small circle. The trap to beware is that this notion no longer agrees with “the number of vertices adjacent to v ” or the “the number of edges incident to v ”

Example 1.2.4

The graph G to the right has two vertices, a and b , and three edges, two between a and b , and a loop at a . Vertex a has degree 4, and vertex b has degree 2.



□

Extended example: Chemistry

In organic chemistry, molecules are frequently drawn as graphs, with the vertices being atoms, and an edge between two vertices if and only if the corresponding atoms have a covalent bond between them (that is, they share a vertex).

Example 1.2.5 Alkanes.

□

The location of an element on the periodic table determines the valency of the element -- hence the degree that vertex has in any molecule containing that graph:

- Hydrogen (H) and Fluorine (F) have degree 1
- Oxygen (O) and Sulfur (S) have degree 2
- Nitrogen (N) and Phosphorous (P) have degree 3
- Carbon (C) has degree 4

Usually, most of the atoms involved are carbon and hydrogen. Carbon atoms are not labelled with a C, but just left blank, while hydrogen atoms are left off completely. One can then complete the full structure of the molecule using the valency of each vertex. On the exam, you may have to know that Carbon has degree 4 and Hydrogen has degree 1; the valency of any other atom would be provided to you

Graphs coming from organic chemistry do not have to be simple – sometimes there are double bonds, where a pair of carbon atoms have two edges between them.

Example 1.2.6 Saturated vs. unsaturated

□

If we know the chemical formula of a molecule, then we know how many vertices of each degree it has. For a general graph, this information is known as the degree sequence

Definition 1.2.7 Degree sequence. The degree sequence of a graph is just the list (with multiplicity) of the degrees of all the vertices. ◇

The following sage code draws a random graph with 7 vertices and 10 edges, and then gives its degree sequence. You can tweak the code to generate graphs with different number of vertices and edges, and run the code multiple times, and the degree sequence should become clear.

```
vertices = 7
edges = 10
```



```

g = graphs.RandomGNM(vertices, edges)
g.show()
print g.degree_sequence()

```

Knowing the chemical composition of a molecule determines the degree sequence of its corresponding graph. However, it is possible that the same set of atoms may be put together into a molecule in more than one different ways. In chemistry, these are called *isomers*. In terms of graphs, this corresponds to different graphs that have the same degree sequence.

An important special case is the constant degree sequence.

Definition 1.2.8 Regular graphs. A graph Γ is *d-regular*, or *regular of degree d* if every vertex $v \in \Gamma$ has the same degree d , i.e. $d(v) = d$. \diamond

As a common special case, a regular graph where every vertex has degree three is called *trivalent*, or *cubic*.

Some quick examples:

1. The cycle graph C_n is two-regular
2. The complete graph K_n is $(n - 1)$ -regular
3. The Petersen graph is trivalent

Handshaking lemma and first applications

To motivate the Handshaking Lemma, we consider the following question. Suppose there seven people at a party. Is it possible that everyone at the party knows exactly three other people?

We can model the situation a graph, with vertices being people at the party, and an edge between two vertices if the corresponding people know each other. The question is then asking for the existence of a graph with seven vertices so that every vertex has degree three. It is then natural to attempt to solve the problem by trying to draw such a graph. After a few foiled attempts, we begin to suspect that it's not possible, but doing a case-by-case elimination of all the possibilities is daunting. It's easier to find a reason why we can't draw such a graph.

We will do this as follows: suppose that everyone at the party who knows each other shakes hands. How many handshakes will occur? On the one hand, from the definitions this would just be the number of edges in the graph. On the other hand, we can count the number of handshakes working person-by-person: each person knows three other people, and so is involved in three handshakes. But each handshake involves two people, and so if we count $7 * 3$ we've counted each handshake twice, and so there should be $7 * 3/2 = 10.5$ handshakes happening, which makes no sense, as we can't have half a handshake. Thus, we have a contradiction, and we conclude such a party isn't possible.

Euler's handshaking Lemma is a generalization of the argument we just made to an arbitrary graph.

Theorem 1.2.9 (*Euler's handshaking Lemma*)

$$\sum_{v \in V(G)} d(v) = 2|E(G)|$$

Proof. We count the “ends” of edges two different ways. On the one hand, every end occurs at a vertex, and at vertex v there are $d(v)$ ends, and so the total number of ends is the sum on the left hand side. On the other hand, every edge has exactly two ends, and so the number of ends is twice the number of

edges, giving the right hand side. ■

We have already seen one use of Euler's handshaking Lemma, but it will be particularly useful in Chapter 3, when we study graphs on surfaces.

Graph Isomorphisms

Generally speaking in mathematics, we say that two objects are "isomorphic" if they are "the same" in terms of whatever structure we happen to be studying. The symmetric group S_3 and the symmetry group of an equilateral triangle D_6 are isomorphic. In this section we briefly discuss isomorphisms of graphs.

Isomorphic graphs

The "same" graph can be drawn in the plane in multiple different ways. For instance, the two graphs below are each the "cube graph", with vertices the 8 corners of a cube, and an edge between two vertices if they're connected by an edge of the cube:




Figure 1.3.1 Two drawings of the cube graph

Example 1.3.2 It is not hard to see that the two graphs above are both drawings of the cube, but for more complicated graphs it can be quite difficult at first glance to tell whether or not two graphs are the same. For instance, there are many ways to draw the Petersen graph that aren't immediately obvious to be the same.

This [animated gif](#) created by Michael Sollami for [this Quanta Magazine article](#) on the Graph Isomorphism problem illustrates many different such drawings in a way that makes the isomorphisms apparent. □

Definition 1.3.3 An isomorphism $\varphi : G \rightarrow H$ of simple graphs is a bijection $\varphi : V(G) \rightarrow V(H)$ between their vertex sets that preserves the number of edges between vertices. In other words, $\varphi(v)$ and $\varphi(w)$ are adjacent in H if and only if v and w are adjacent in G . ◇

Example 1.3.4




Figure 1.3.5 C_5 is isomorphic to its complement C_5^c

The cycle graph on 5 vertices, C_5 is isomorphic to its complement, C_5^c . The cycle C_5 is usually drawn as a pentagon, and if we were then going to naively draw C_5^c we would draw a 5-sided star. However, we could draw C_5^c differently as shown, making it clear that it is isomorphic to C_5 , with isomorphism $\varphi : C_5 \rightarrow C_5^c$ defined by $\varphi(a) = a, \varphi(b) = c, \varphi(c) = e, \varphi(d) = b, \varphi(e) = d$. \square

Although solving the graph isomorphism problem for general graphs is quite difficult, doing it for small graphs by hand is not too bad and is something you must be able to do for the exam. If the two graphs are actually isomorphic, then you should show this by exhibiting an isomorphism; that is, writing down an explicit bijection between their vertex sets with the desired properties. The most attractive way of doing this, for humans, is to label the vertices of both copies with the same letter set.

If two graphs are not isomorphic, then you have to be able to prove that they aren't. Of course, one can do this by exhaustively describing the possibilities, but usually it's easier to do this by giving an obstruction – something that is different between the two graphs. One easy example is that isomorphic graphs have to have the same number of edges and vertices. We'll discuss some others in the next section

Heuristics for showing graphs are or aren't isomorphic

Another, only slightly more advanced invariant is the degree sequence of a graph that we saw last lecture in our discussion of chemistry.

If $\varphi : G \rightarrow H$ is an isomorphism of graphs, then we must have $d(\varphi(v)) = d(v)$ for all vertices $v \in G$, and since isomorphisms are bijections on the vertex set, we see the degree sequence must be preserved. However, just because two graphs have the same degree sequences does not mean they are isomorphic.

Slightly subtler invariants are number and length of cycles and paths.

Cultural Literacy: The Graph Isomorphism Problem

This section, as all "Cultural Literacy" sections, is information that you may find interesting, but won't be examined.

The graph isomorphism problem is the following: given two graphs G and H , determine whether or not G and H are isomorphic. Clearly, for any two graphs G and H , the problem is solvable: if G and H both of n vertices, then there are $n!$ different bijections between their vertex sets. One could simply examine each vertex bijection in turn, checking whether or not it maps edges to edges.

The problem is interesting because the naive algorithm discussed above is not very efficient: for large n , $n!$ is absolutely huge, and so in general this

algorithm will take a long time. The question is, is there a faster way to do check? How fast can we get?

The isomorphism problem is of fundamental importance to theoretical computer science. Apart from its practical applications, the exact difficulty of the problem is unknown. Clearly, if the graphs are isomorphic, this fact can be easily demonstrated and checked, which means the Graph Isomorphism is in NP.

Most problems in NP are known either to be easy (solvable in polynomial time, P), or at least as difficult as any other problem in NP (NP complete). This is not true of the Graph Isomorphism problem. In November of last year, Laszlo Babai announced a quasipolynomial-time algorithm for the graph isomorphism problem – you can read about this work in this great popular science article.

Instant Insanity

So far, our motivation for studying graph theory has largely been one of philosophy and language. Before we get too much deeper, however, it may be useful to present a nontrivial and perhaps unexpected application of graph theory; an example where graph theory helps us to do something that would be difficult or impossible to do without it.

A puzzle



Figure 1.4.1 Instant Insanity Package

There is a puzzle marketed under the name "Instant Insanity", one version of which is shown above. The puzzle is sometimes called "the four cubes problem", as it consists of four different cubes. Each face of each cube is painted one of four different colours: blue, green, red or yellow. The goal of the puzzle is to line the four cubes up in a row, so that along the four long edges (front, top, back, bottom) each of the four colours appears exactly once.

Depending on how the cubes are coloured, this may be not be possible, or there may be many such possibilities. In the original instant insanity, there is exactly one solution (up to certain equivalences of cube positions). The point of the cubes is that there are a large number of possible cube configurations, and so if you just look for a solution without being extremely systematic, it is highly unlikely you will find it.

But trying to be systematic and logical about the search directly is quite difficult, perhaps because we have problems holding the picture of the cube in

our mind. In what follows, we will introduce a way to translate the instant insanity puzzle into a question in graph theory. This is obviously in no way necessary to solve the puzzle, but does make it much easier. It also demonstrates the real power of graph theory as a visualization and thought aid.

There are many variations on Instant Insanity, discussions of which can be found [here](#) and [here](#). There's also a [commercial](#) for the game.

Enter graph theory

It turns out that the important factor of the cubes is what color is on the opposite side of each face. Suppose we want face one facing forward. Then we have four different ways to rotate the cube to keep this the same. The same face will always appear on the opposite side, but we can get any of the remaining four faces to be on top, say.




Figure 1.4.2 An impossible set of cubes

Let us encode this information in a graph. The vertices of the graph will be the four colors, B, G, R and Y. We will put an edge between two colors each time they appear as opposite faces on a cube, and we will label that edge with a number 1-4 denoting which cube the two opposite faces appear. Thus, in the end the graph will have twelve edges, three with each label 1-4. For from the first cube, there will be a loop at B, and edge between G and R, and an edge between Y and R. The graph corresponding to the four cubes above is the following:




Figure 1.4.3 The graph constructed from the cubes in [Figure 1.4.2](#)

Proving that our cubes were impossible

We now analyze the graph to prove that this set of cubes is not possible.

Suppose we had an arrangement of the cubes that was a solution. Then, from each cube, pick the edge representing the colors facing front and back on that cube. These four edges are a subgraph of our original graph, with one edge of each label, since we picked one edge from each cube. Furthermore, since we assumed the arrangement of cubes was a solution of instant insanity, each color appears once on the front face and once on the back. In terms of our subgraph, this translates into asking that each vertex has degree two.

We can get another subgraph satisfying these two properties by looking at the faces on the top and bottom for each cube and taking the corresponding edges. Furthermore, these two subgraphs do not have any edges in common.

Thus, given a solution to the instant insanity problem, we found a pair of subgraphs S_1, S_2 satisfying:

1. Each subgraph S_i has one edge with each label 1,2,3,4
2. Every vertex of S_i has degree 2
3. No edge of the original graph is used in both S_1 and S_2

As an exercise, one can check that given a pair of subgraphs satisfying 1-3, one can produce a solution to the instant insanity puzzle.

Thus, to show the set of cubes we are currently examining does not have a solution, we need to show that the graph does not have two subgraphs satisfying properties 1-3.

To do, this, we catalog all graphs satisfying properties 1-2. If every vertex has degree 2, either:

1. Every vertex has a loop
2. There is one vertex with a loop, and the rest are in a triangle
3. There are two vertices with loops and a double edge between the other two vertices

4. There are two pairs of double edges
5. All the vertices live in one four cycle
6. A subgraphs of type 1 is not possible, because G and R do not have loops.

For subgraphs of type 2, the only triangle is G-R-Y, and B does have loops. The edge between Y-G must be labeled 3, which means the loop at B must be labeled 1. This means the edge between G and R must be labeled 4 and the edge between Y and R must be 2, giving the following subgraph:




Figure 1.4.4 A subgraph for a solution for one pair of faces

For type 3, the only option is to have loops at B and Y and a double edge between G and R. We see the loop at Y must be labeled 2, one of the edges between G and R must be 4, and the loop at B and the other edge between G and R can switch between 1 and 3, giving two possibilities:




Figure 1.4.5 Two more subgraphs for a partial solutions



For subgraphs of type 4, the only option would be to have a double edge between B and G and another between Y and R; however, none of these edges are labeled 3 and this option is not possible.

Finally, subgraphs of type 5 cannot happen because B is only adjacent to G and to itself; to be in a four cycle it would have to be adjacent to two vertices that aren't itself.

This gives three different possibilities for the subgraphs SiSi that satisfy properties 1 and 2. However, all three possibilities contain the the edge labeled 4 between G and R; hence we cannot choice two of them with disjoint edges, and the instant insanity puzzle with these cubes does not have a solution.

Other cube sets

The methods above also give a way to find the solution to a set of instant insanity cubes should one exist. I illustrate this in the following Youtube video:



YouTube: <https://www.youtube.com/watch?v=GsbhRfjaaN8>

Trees

A very important class of graphs are trees -- they are connected, but just barely: removing any edge causes them not to be connected.

Basics on trees




Figure 1.5.1 A forest consisting of three trees

The figure above shows some examples of the trees. Meanwhile, the cycle graph C_n or the complete graph K_n with $n \geq 3$ are not trees: we can remove an edge from these graphs and they'd still be connected. The formal definition of a tree is as follows:

Definition 1.5.2 Trees and Forests. A graph G is a *tree* if:

1. G is connected
2. G has no cycles

A not necessarily connected graph with no cycles is called a *forest*, so that a forest is a union of trees. \diamond

Informally, the condition that G is connected is asking that G have enough edges, while the condition that G has no cycles is asking for G not to have too many edges. Thus, trees are sort of "goldilocks" graphs -- they have enough edges, but not too many -- they're connected, but just barely.

The following Theorem gives many alternate characterisations of trees, and makes more precise the intuition of trees as "goldilocks graphs".

Theorem 1.5.3 Let G be a graph with n vertices. The following are equivalent:

1. G is a tree.
2. Between any two vertices $a, b \in V(G)$, there is a unique path.
3. G is connected, but removing any edge makes G disconnected.
4. G has no cycles, but adding any edges to G creates a cycle.
5. G is connected and has $n - 1$ edges
6. G has no cycles and has $n - 1$ edges

We will not use most of , and will not prove that all options are equivalence. We briefly proof that 1 is equivalent to 2 now, and in the next section we will carefully prove that 1 is equivalent to 5, as we will use this fact a few times. The rest make a good exercise to check your basic understanding.

To see that 1 and 2 are equivalent, note that being connected by definition means there is a path between every two vertices. As a tree is a connected graph without any cycles, to finish seeing 1 and 2 are equivalent is exactly [Lemma 1.5.5](#), whose main idea is contained in [Figure 1.5.4](#)




Figure 1.5.4 Two different paths create a cycle

Lemma 1.5.5 *A graph G has no cycles if and only if there is at most one path between any two vertices of G .*

Proof. If G has a cycle, then there are at least two paths between any vertex on that cycle -- the paths going each way around the cycle. Thus, we just have to show that if there are two paths between v and w , then G has a cycle.

In the easy case, the two paths will contain no vertices in common except for v and w , and so the union of the two paths will be a cycle. In general, the paths will share other vertices and edges -- they may well repeat vertices and edges themselves. But in any case, by considering some subset of these two paths will be a cycle. ■

Leaves

To prove Theorem , we first need to introduce the concept of leaves.

Definition 1.5.6 Leaf. A vertex $v \in G$ is called a *leaf* if it has degree one, i.e. if $d(v) = 1$ ◇

When looked at a drawn graph, this definition is fairly intuitive: real life trees branch out and split in leaves, just like mathematical trees.

Lemma 1.5.7 Trees have leaves. *Let T be a finite tree with at least two vertices. Then T has at least two leaves.*

Proof. By assumption, T has at least two vertices, say v_0 and w_0 . Since T is a tree it is connected, and so in particular there must be a path between v_0 and w_0 ; let v_i be the vertices in this path, and let e_i be the edge in the path joining v_{i-1} to v_i .

Since v_m is adjacent to e_m it has degree at least one; if it has degree 1 it is a leaf, and we've found a leaf. If v_m is not a leaf, then there must be another edge coming out of it, say e_{m+1} going to v_{m+1} . Note that v_{m+1} cannot be any of the vertices we've already found, as then we'd have more than one path between two vertices and hence a loop, but T was a tree. Thus we can make the path a bit longer.

We can now continue this argument inductively as long as the vertex at the end of the path has degree higher than 1. But since T is finite, and we never return to a vertex we've already visited, this process must eventually terminate, but the only way this can happen is if the end vertex of the path has degree 1, that is, if it's a leaf.

A similar argument extending from v_0 , the other end of the path, shows

that we must eventually reach a different leaf from that end, and so **T** must have at least two leaves, as desired. ■

Now that we have a basic understanding of leaves, we are ready to prove the following:

Lemma 1.5.8 *A simple graph **G** with n vertices is a tree if and only if is connected and has $n - 1$ edges.*

Proof. Since being connected is half of the requirement of being a tree, we need to show that a connected graph on n vertices is a tree if and only if it has $n - 1$ edges.

Let us first prove that a tree on n vertices has $n - 1$ edges. We will use induction on n . As bases, there is only one tree with 1 vertex, and it does in fact have 0 edges, and there is only one tree with two vertices, and it does in fact have 1 edge. So for the inductive step, let us suppose that we know that all trees with $k < n$ vertices have $k - 1$ edges, and let **T** be a tree with n vertices. By Lemma, we know that **T** has a leaf v , which by the definition of leaves is connected to the rest of **T** by a single edge e . If we delete v and e from **T**, we get a smaller graph **T'** in which has one less vertex and one less edge than **T** did.

Since **T** was a tree, it follows that **T'** is a tree, too -- check this yourself, using the definition of a tree! Then, since **T'** is a tree with $n - 1$ vertices, by the inductive hypothesis it follows that it has one less edge $n - 2$ edges, and so **T** must have $n - 1$ edges, which is what we were trying to show.

Now we show the other direction: that if **G** is a simple connected graph with n vertices and $n - 1$ edges, then **G** is a tree. The basic structure of the proof is the same as the other direction: find a vertex v adjacent to a single edge e , and delete v and e to get a smaller tree, where we can assume the proposition holds, and then use induction.

To play the role of the lemma that every tree has a leaf, we will establish the following statement: if **G** is a connected graph with n vertices and $n - 1$ edges, then **G** has a vertex of degree 1. Note that since **G** is connected, it can't have any vertices of degree 0, and so to prove it has a vertex of degree 1 it is enough to show that it has a vertex of degree strictly less than 2. Thus, for contradiction assume that every vertex v of **G** has degree $d(v) \geq 2$. But then the handshaking lemma would say:

$$2(n - 1) = 2e = \sum d(v) \geq \sum 2 = 2n$$

a contradiction, and hence **G** must have a vertex of degree 1, as desired. ■

Trees in Chemistry

Our brief study of trees has the following consequence for chemistry: whether or not a molecule is a tree is determined just by its chemical formula, and not how it's put together. Equivalently, if one isomer of a molecule is a tree, then all isomers of the molecule are.

The argument runs as follows. Knowing the chemical formula of a molecule means we know the degree sequence of the corresponding graphs. Molecules are by definition connected graphs, so to be a tree it is enough to show that the graph has one less edge than the number of vertices. But we can compute the number of edges from the degree sequence by using the Handshaking Lemma.

Example 1.5.9 Alkanes are trees. Any molecule with formula C_nH_{2n+2} is an *alkane*. Although in general alkanes can have multiple isomers, every isomer of an Alkane will always be a tree, as we now show.

To show a graph is a tree, it suffices to show that it is connected and that the number of edges is one less than the number of vertices. Since Alkanes are molecules, we know that the graph is connected. Furthermore, C_nH_{2n+2} has n carbons and $2n + 2$ hydrogens, and hence has $3n + 2$ vertices. Thus, it is enough to show that any molecule with formula C_nH_{2n+2} has $3n + 1$ edges.

To do this, we use the handshaking lemma: $2e = \sum d(v)$. Each of the n carbons has degree 4, so the carbons contribute $4n$ to the total degree, and each hydrogen of the $2n + 2$ has degree 1 and so only contributes 1 to the sum of degrees. Hence, twice the number of edges is equal to $4n + (2n + 2) = 6n + 2$, and so there are $3n + 1$ edges, as desired. \square

As an early application of graph theory, Cayley used these ideas to count the number of isomers of Alkanes (with some mistakes). In general, you can count isomers of any molecule by counting isomorphism classes of graphs with given degree sequences, but it can help organize the search to know, e.g., that they're all trees. To make sure we don't miss or double count any, it's useful to organize the enumeration according to some principle; for Alkanes one way is to organize according to the longest path of carbons, another is to restrict degree sequences to just how the carbons have connected to other carbons.

Example 1.5.10 Counting isomers of C_6H_{14} . We illustrate these both of these methods. We first illustrate by length of the longest path of carbons.

- Chain length six: Since we've used all carbons then there are no more choices, and there is only one such isomer.
- Chain length five: We need to add one more carbon. We can't add it to either of the end carbons, because then we'd have a path of length 6. We can add it to the central of the two chains, or alternatively to one either side of central -- the last two trees are isomorphic, giving us two possibilities
- Chain length four: We need to add two more carbons. Again, they can't be added to either of the end carbons, or we'd have a longer chain length. Therefore, they must be added to the two central carbons. One case is that we add one carbon to each of the two central carbons. Alternatively, we could add both the carbons to the same "central" carbon reversing the order of the chain is a symmetry that interchanges the two central atoms. We could add each carbon directly to the existing carbon in the chain, or we could add them one after the other making a chain of length two. However, the resulting graph would have a chain of length 5 and already be counted. Thus, there are two possibilities here.
- Chain length three: We need to add three more carbons, and there's only one central carbon to attach them to. We can't add them all directly to this central carbon, as that would create a carbon of degree 5. On the other hand, once we add a chain of length longer than one to this central carbon we'd have a path of length 4 or greater.

Thus, we see there are five isomers of C_6H_{14} . Alternatively, we could organize our search by deleting the hydrogens, and then considering the degrees of the resulting carbon-carbon graph.

- Degree at most two: If the resulting tree only had carbons of degree at most two, then it would have to be the path graph P_6 , and so we only have one possibility here.
- One vertex of degree three: If the resulting graph had exactly one carbon of degree three, that vertex and its three neighbours would account for

4 of our 6 carbons, and so we'd have to add two more. We couldn't add them directly to the same vertex, as that would create a second vertex of degree three. So, they could either be added as a chain of length two to one of the leaves of the existing graph, or they could be added to two separate leaves. Drawing these graphs we see they're not isomorphic, and so we have two possibilities here.

- Two vertices of degree three: If we have two vertices of degree three, one sees they'd have to be adjacent to each other, resulting in one possibility.
- Vertex of degree four: A vertex of degree 4 and its four neighbours would account for all but one of the carbons. We could add that carbon to any of the leaves, and get one more possibility.




Figure 1.5.11 Isomers of C_6H_{12}

Since carbons only have degree 4, the tree with six vertices where all are connected to a central vertex isn't allowed, and we have found all the isomers. \square

Exercises

- For each of the following sequences, either give an example of such a graph, or explain why one does not exist.
 - A graph with six vertices whose degree sequence is $[5, 5, 4, 3, 2, 2]$
 - A graph with six vertices whose degree sequence is $[5, 5, 4, 3, 3, 2]$
 - A graph with six vertices whose degree sequence is $[5, 5, 5, 5, 3, 3]$
 - A simple graph with six vertices whose degree sequence is $[5, 5, 5, 5, 3, 3]$
- For the next Olympic Winter Games, the organizers wish to expand the number of teams competing in curling. They wish to have 14 teams enter, divided into two pools of seven teams each. Right now, they're thinking of requiring that in preliminary play each team will play seven games against distinct opponents. Five of the opponents will come from their own pool and two of the opponents will come from the other pool. They're having trouble setting up such a schedule, so they've come to you. By using an appropriate graph-theoretic model, either argue that they cannot use their current plan or devise a way for them to do so.

3. Figure 1.6.1 contains four graphs on six vertices. Determine which (if any) pairs of graphs are isomorphic. For pairs that are isomorphic, give an isomorphism between the two graphs. For pairs that are not isomorphic, explain why.




Figure 1.6.1 Are these graphs isomorphic?

4. Let G be a simple graph with n vertices and degree sequence a_1, a_2, \dots, a_n . What's the degree sequence of its complement G^c ?
5. Let G be the graph with vertices consisting of the 10 three element subsets of $\{a, b, c, d, e\}$, and two vertices adjacent if they share exactly one element. So, for example, the two vertices $v = \{a, c, e\}$ and $w = \{b, c, d\}$ are adjacent, but neither v or w is adjacent to $u = \{a, b, c\}$.
 Draw G in a way that shows it is isomorphic to the Petersen graph.
 Now let H be the graph with vertices consisting of the 10 two element subsets of $\{a, b, c, d, e\}$, and two vertices adjacent if they share *no* elements. Without drawing H , write down an isomorphism between G and H . Hint: There's a "natural" bijection between the two and three element subsets of $\{a, b, c, d, e\}$.
6. Recall that G^c denotes the complement of a graph G . Prove that $f : G \rightarrow H$ is an isomorphism of graphs if and only if $f : G^c \rightarrow H^c$ is an isomorphism.
7. Determine the number of non-isomorphic simple graphs with seven vertices such that each vertex has degree at least five.
Hint. Consider the previous exercise
8. Consider the standard Instant Insanity puzzle, with four cubes and four colours. Explain why one would expect there to be 331,776 different cube configurations. Further explain why there would be fewer configurations if any cubes are coloured with symmetries.
 In the text, we solve the puzzle by finding certain pairs of subgraphs. Assuming that none of the cubes are coloured symmetrically, explain why each pair of subgraphs corresponds to at least 8 different cube configurations that are actually solutions, and why, depending on the isomorphism type of the subgraphs found, there may be more solutions.
9. Variations of the Instant Insanity puzzle increase the number of cubes and the number of colours. Explain how to modify our graph theoretic solution to solve the puzzle when we have n cubes, each face of which is coloured one of n colours, and we want to line up the cubes so that each of the top, bottom, front and rear strings of cubes displays each of the n colours

exactly once.

10. Use the method from the previous question to solve the following set of six cubes, marketed under the name "Drive ya crazy", where each face is coloured either blue, cyan, green, orange, red, or yellow.




Figure 1.6.2 The six cubes from "Drive Ya crazy"

Chapter 2

Walks

In this chapter we investigate walks in graphs. We first look at some basic definitions and examples, we discuss Dijkstra's algorithm for finding the shortest path between two points in a weighted graph, and we discuss the notions of Eulerian and Hamiltonian graphs.

Walks: the basics

If the edges in a graph Γ represent connections between different cities, it is obvious to start planning longer trips that compose several of these connections. The notion of a *walk* formally captures this definition; the formal notions of *path* and *trail* further ask that we not double back on ourselves or repeat ourselves in certain formally defined ways.

Once we've done that, we investigate what it means for a graph to be connected or disconnected.

Walks and connectedness

Before we see the formal definition of a walk, it will be useful to see an example:




Figure 2.1.1 Example of a walk

In the graph shown, the vertices are labelled with letters, and the edges are labelled with numbers, and we have a walk highlighted in red, and with arrowtips drawn on the edges. Starting from vertex A , we can take edge 6 to vertex D , and then edge 4 to vertex C , edge 5 to vertex F , edge 3 back to vertex D , and finally edge 8 to vertex E . Intuitively, then, a walk strings together several edges that share vertices in between. Making that formal, we have the following.

Definition 2.1.2 Walk. *walk* in a graph Γ is a sequence

$$v_0, e_1, v_1, e_2, v_2, \dots, v_{n-1}, e_n, v_n$$

where the v_i are vertices, the e_j are edges, and the edge e_j goes between vertices v_{j-1} and v_j .

We say that the walk is between vertices $a = v_0$ and $b = v_n$ \diamond

With this notation for a walk, Example [Figure 2.1.1](#), the walk shown would be written $A, 6, D, 4, C, 5, F, 3, D, 8, E$. The visual representation of the walk on the graph is vastly more intuitive, the written one feeling cumbersome in comparison.

The definition of walk above contains some extra information. If we just know the sequence of edges we can reconstruct what the vertices have to be (assuming we have at least two edges in the walk). Alternatively, if the graph Γ does not have multiple edges, it is enough to just know the vertices v_i , but if Γ has multiple edges that just knowing the vertices does not determine the walk.

Definition 2.1.3 Connected. We say a graph Γ is *connected* if for any two vertices v, w , there is a walk from v to w in Γ . \diamond

Definition 2.1.4 Disjoint union. Given two graphs Γ_1 and Γ_2 , the *disjoint union* $\Gamma_1 \sqcup \Gamma_2$ is obtained by taking the disjoint union of both the vertices and edges of Γ_1 and Γ_2 . So $\Gamma_1 \sqcup \Gamma_2$ consists of a copy of Γ_1 and a copy of Γ_2 , with

no edges in between the two graphs. \diamond

Definition 2.1.5 Disconnected. A graph Γ is *disconnected* if we can write $\Gamma = \Gamma_1 \sqcup \Gamma_2$ for two proper (i.e., not all of Γ) subgraphs Γ_1 and Γ_2 . \diamond

We now have a definition for what it means for a graph to be connected, and another for what it means for a graph to be disconnected. From everyday usage of this words, we would certainly hope that a graph is disconnected if and only if it is not connected. However, it is not immediately clear from the definitions as written that this is the case.

Lemma 2.1.6 *The following are equivalent:*

1. Γ is connected
2. Γ is not disconnected

Proof. 1 implies 2: Suppose that Γ is connected, and let $v, w \in V(\Gamma)$; we want to show that there is a walk from v to w .

Define $S \subset V(\Gamma)$ to be the set of all vertices $u \in V(\Gamma)$ so that there is a walk from v to u ; we want to show that $w \in S$.

First, observe that there are no edges from S to $V(\Gamma) \setminus S$. Suppose that e was an edge between $a \in S$ and $b \in V(\Gamma) \setminus S$. Since $a \in S$, by the definition of S there is a walk $v = v_0 v_1 v_2 \cdots v_m = a$ from v to a . We can add the edge e to the end of the walk, to get a walk from v to b , and hence by definition $b \in S$.

Now suppose that $w \notin S$. Then S and $V(\Gamma) \setminus S$ are both nonempty, and by the above there are no edges between them, and so Γ is not connected, a contradiction.

To prove 2 implies 1, we prove the contrapositive. If Γ is not connected, then there are two vertices $v, w \in V(\Gamma)$ so that there is no walk from v to w .

Suppose that $\Gamma = \Gamma_1 \sqcup \Gamma_2$, and pick $v \in V(\Gamma_1), w \in V(\Gamma_2)$. Any walk from v to w starts in $V(\Gamma_1)$ and ends in $V(\Gamma_2)$, and so at some point there must be an edge from a vertex in Γ_1 to a vertex in Γ_2 , but there are no such edges \square

Types of Walks

Many questions in graph theory ask whether or not a walk of a certain type exists on a graph: we introduce some notation that will be needed for these questions.

Definition 2.1.7 We say a walk is *closed* if it starts and ends on the same vertex; i.e., $v_0 = v_n$. The *length* of a walk is n , the number of edges in it. The *distance* between two vertices v and w is the length of the shortest walk from v to w , if one exists, and ∞ if one does not exist. \diamond

It is sometimes convenient to have terminology for walks that don't back-track on themselves:

Definition 2.1.8

1. If the edges e_i of the walk are all distinct, we call it a *trail*
2. If the vertices v_i of the walk are all distinct (except possibly $v_0 = v_m$), we call the walk a *path*. The exception is to allow for the possibility of closed paths.

\diamond

Lemma 2.1.9 *Let $v, w \in V(\Gamma)$. The following are equivalent:*

1. There is a walk from v to w
2. There is a trail from v to w

3. *There is a path from v to w .*

As is often the case, the formal write-up of the proof makes something that can seem very easy intuitively look laborious, so it's worth analysing it briefly for our example walk $A - D - C - F - D - E$ from Figure 2.1.1. This walk is not a path as it repeats the vertex D ; however, we may simply remove the triangle $D - C - F - D$ from the walk to get the trail $A - D - E$. this idea is what works in general.

Proof. It is immediate from the definitions that 3 implies 2 which implies 1, as any path is a trail, and any trail is a walk.

That 1 implies 3 is intuitively obvious: if you repeat a vertex, then you've visited someplace twice, and weren't taking the shortest route. Let's make this argument mathematically precise.

Suppose we have a walk $v = v_0, e_1, \dots, e_m, v_m = w$ that is not a path. Then, we must repeat some vertex, say $v_i = v_k$, with $i < k$. Then we can cut out all the vertices and edges between v_i and v_k to obtain a new walk

$$v = v_0, e_1, v_1, \dots, e_i, v_i = v_k, e_{k+1}, v_{k+1}, e_{k+2}, v_{k+2}, \dots, v_m$$

Since $i < k$, the new walk is strictly shorter than our original walk. Since the length of a walk is finite, if we iterate this process the result must eventually terminate. That means all our vertices are distinct, and hence is a path. ■

Eulerian Walks

In this section we introduce the problem of Eulerian walks, often hailed as the origins of graph theory. We will see that determining whether or not a walk has an Eulerian circuit will turn out to be easy; in contrast, the problem of determining whether or not one has a Hamiltonian walk, which seems very similar, will turn out to be very difficult.

The bridges of Konigsburg

The city of Konigsburg (now Kaliningrad) was built on two sides of a river, near the site of two large islands. The four sectors of the city were connected by seven bridges, as follows (picture from Wikipedia):

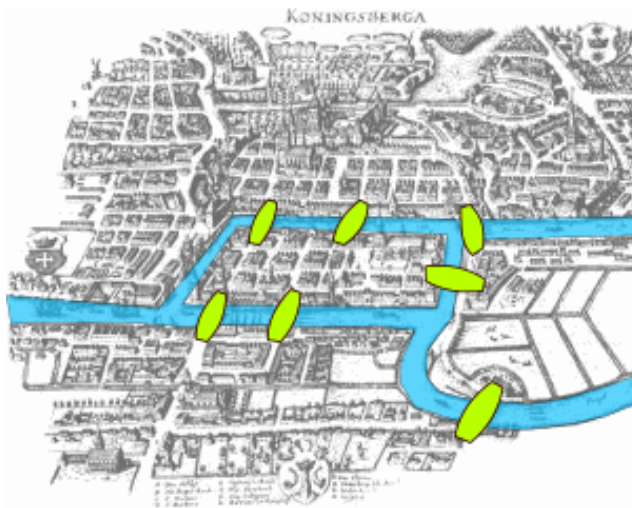


Figure 2.2.1 The city of Konigsburg in Euler's time

A group of friends enjoyed strolling through the city, and created a game: could they take a walk in the city, crossing every bridge exactly once, and return to where they started from? They couldn't find such a walk, but they couldn't prove such a walk wasn't possible, and so they wrote to the mathematician Euler, who proved that such a walk is not possible.

Eulerian Walks: definitions

We will formalize the problem presented by the citizens of Königsburg in graph theory, which will immediately present an obvious generalization.

We may represent the city of Königsburg as a graph Γ_K ; the four sectors of town will be the vertices of Γ_K , and edges between vertices will represent the bridges (hence, this will not be a simple graph).

Then, the question reduces to finding a closed walk in the graph that will use every edge *exactly* once. In particular, this walk will not use any edge more than once and hence will be a trail.

Checkpoint 2.2.2 Which graphs have walks that use every vertex and edge exactly once?

Definition 2.2.3 Let G be a graph. An *Eulerian cycle* is a closed walk that uses every edge of G exactly once.

If G has an Eulerian cycle, we say that G is *Eulerian*.

If we weaken the requirement, and do not require the walk to be closed, we call it an Euler path, and if a graph G has an Eulerian path but not an Eulerian cycle, we say G is *semi-Eulerian*. \diamond

The question of the walkers of Königsburg is then equivalent to asking if the graph Γ_K is Eulerian. The birth of graph theory is usually marked to the following theorem, proven by Euler:

Theorem 2.2.4 A connected graph Γ is Eulerian if and only if every vertex of Γ has even degree

A digression on proofs, formality, and intuition

Before discussing the proof of [Theorem 2.2.4](#), it's worth a little meta-discussion about proofs, intuition vs. rigor, and mathematics as a whole. The proofing [Theorem 2.2.4](#) is a common exam question, and you may not be used to studying for reproducing proofs on exams. Certainly one way to prepare for such a question is to memorize the proof word for word. There doesn't seem to be a lot of obvious value in this approach, however. So why ask these questions on the exam? And this opens the door to more philosophical questions as well: how should we think/interact with proofs anyway? What's the point of it all?

Usually in books or in lectures, proofs are only given in slick, elegant, polished formal versions. There are many reasons for this: there's a certain beauty to it; it's important to write it out formally to make sure it's all correct; there's only so much time in lectures, and brevity is a virtue anyway. People turn away from long works, and the main point of a proof, after all, is to prove something, and it's easier to check that it's all correct if it's shorter.

But there's a very real downside to this presentation of proofs as the finished, elegant thing. Most important to me is that the way mathematics is written formally on the page is very different from how it lives actively in our brains (or my brain, at least). Nobody (or certainly very few people) comes up with proofs in the elegant short start to finish way that they're written. Typically, there's a mess of chaotic half ideas that slowly get refined down to the written proof that you see. But often the mess is the exciting part,

We sketch a few of the main ideas of the proof in an informal setting now, before giving a complete formal proof. To learn this proof for the exam, you should have this informal picture in your head, and perhaps a skeleton outline of the main formal points that need to be shown. You shouldn't try to memorize the formal proof word for word like a poem; instead, practice expanding out from the informal ideas/skeleton proof to the full formal proof on your own a few times.

It is much easier to see that if a graph G is Eulerian, then every vertex has even degree. Suppose we wanted to show that a given vertex v was Eulerian; let us stand at the vertex v and have a friend trace out the Eulerian cycle. We'll wait for a while, and then the friend will appear at v along some edge e_1 , and then live along some different edge e_2 . We'll wait some more, and they'll reappear coming from new edge e_3 , and then leave again along some edge e_4 .

This will continue until they have arrived or left by every edge that hits v . But every time they visit v , they must arrive by one edge, and leave by another one, and hence every visit uses up an even number of edges, and so the degree $d(v)$ of v must be even. But there was nothing special about the vertex v , and hence the degree of every vertex must be even.

To argue the other way is more difficult; before trying to show there's a closed path that uses all the edges, let's just construct *any* closed path. We pick some vertex v_0 to start at, and just randomly choose an edge out of v_0 , to some other vertex v_1 , and from there randomly choosing any edge we haven't used yet to another vertex v_n , and so on.

To construct a closed walk, we'd like to show we eventually have to return to v_0 . We're only working with finite graphs, so our walk can't continue forever; the only possibility we have to rule out is that we reach some vertex v_n and find that we have already used every vertex incident to v_n . But as we saw before, the path will pair up the edges incident to each vertex as an arriving edge and a departing edge, and we know the degree of v_n is even. If the path has already visited $v_n k$ times, then we'll have used $2k$ of the edges incident to it; when we arrive for the $k + 1$ st time we'll use one edge, and in all we'll have used $2k + 1$ edges, an odd number; since the degree of v_n is even there must be at least one edge we haven't used to exit by.

You might worry that the argument above suggests we can carry on the walk forever, which we obviously can't do since the graph is finite, but the argument above doesn't work for v_0 : when we start the path out at v_0 , we haven't had to arrive there, and so the edge we initially leave by is not paired with anything. Therefore, if Γ has all vertices with even degree, and we walk randomly for as long as we can, we'll always get stuck at our starting vertex.

Hence, we have shown that if Γ is a graph with all vertices having even degree, there will exist some closed walk in Γ , but the walk we created was chosen randomly, and there's no guarantee it will include all the edge of Γ -- in all likelihood, it won't.

But if we look at the edges we haven't used, they will form a simpler graph. Γ' . There's no reason for Γ' to be connected, but it's not too hard to see that every vertex of Γ' will still have even degree: in Γ every vertex had even degree, and we saw in our first proof that a closed walk that doesn't repeat edges uses up an even number of edges at each vertex, and so we'll have an even number of edges left at each point. Thus, each connected piece of Γ' satisfies the hypothesis of the problem, and is simpler, so we can try to find a closed walk on each of connected piece of Γ' , and then "stitch" the results together to get a walk that uses all the pieces. In the formal proof, this process is best captured using induction, and we'll save the complete description until then, but for now we illustrate the process in an example




Figure 2.2.5 Example of a walk

Example 2.2.6 Let's see an example of how the process of finding an Eulerian path works for the graph Γ in Figure below.




Figure 2.2.7 A graph Γ

It probably isn't hard to immediately find an Eulerian cycle for Γ just by examination, but to illustrate the algorithm to begin with, we are going to deliberately choose a cycle that doesn't use every edge, the cycle `aeijgba` show in the next figure:




Figure 2.2.8 Initializing with a closed walk in Γ

To extend our cycle to an Eulerian cycle, we delete all the edges used in the graph, and study the remaining graph. In our case, this graph has two connected components, a four cycle and a three cycle. Both of these are cycles themselves, they're trivially Eulerian, -- in general, it might take some work to find an Eulerian cycle for the components, but won't be too hard as the graph will be smaller.




Figure 2.2.9 Parts of Γ missed by our walk

Finally, we stitch our Eulerian paths together. We follow our initial cycle that wasn't an Eulerian cycle, and the first time we hit a vertex that's in one of the other cycles, we insert that cycle in before we continue along our original path. In our example, our original cycle was $aeijgba$. a is not either of the other components, but e is in the four cycle, so before we continue on our original cycle to i , we insert the four cycle $ehie$, giving $aeihje$. We now continue along our original cycle, adding ei , ij , ig , until we reach a vertex g that's in one of the added cycles, which we then insert, giving $aeihjeijgcdg$ as our cycle so far.

Continuing this process, we find $aeihjeijgcdgba$ as an Eulerian cycle.




Figure 2.2.10 The Eulerian cycle stitched together

Proof.

□
■

Remark 2.2.11 Note that it does *not* say: "A graph Γ is Eulerian if and only if it is connected and every vertex has even degree." This statement in quotation marks is false, but for "stupid" reasons. If Γ is Eulerian, and E_n is the graph with n vertices with no edges, then $\Gamma \sqcup E_n$ is Eulerian but not connected. These are the only examples of such graphs.

Theorem 2.2.12 A connected graph Γ is semi-Eulerian if and only if it has exactly two vertices with odd degree.

Proof. A minor modification of our argument for Eulerian graphs shows that the condition is necessary. Suppose that Γ is semi-Eulerian, with Eulerian path $v_0, e_1, v_1, e_2, v_3, \dots, e_n, v_n$. Then at any vertex other than the starting or ending vertices, we can pair the entering and leaving edges up to get an even number of edges.

However, at the first vertex v_0 the path leaves along e_1 the first time but never enters it accordingly, so that v_0 has an odd degree; similarly, at v_n the path enters one final time along e_n without leaving, and so v_n also has an odd degree.

To see the condition is sufficient we could also modify the argument for the Eulerian case slightly, but it is slicker instead to *reduce* to the Eulerian case. Suppose that Γ is connected, and that vertices v and w have odd degree and all other vertices of Γ have even degree. Then we can construct a new graph Γ' by adding an extra edge $e = vw$ to Γ . Then Γ' is connected and every vertex has even degree, and so it has an Eulerian cycle. Deleting the edge e that we added from this cycle gives an Eulerian path from v to w in Γ . ■

Hamiltonian cycles

We now introduce the concept of Hamiltonian walks. Though on the surface the question seems very similar to determining whether or not a graph is Eulerian, it turns out to be much more difficult.

Definition 2.3.1 A graph is *Hamiltonian* if it has a closed walk that uses every vertex exactly once; such a path is called a *Hamiltonian cycle* ◇

First, some very basic examples:

1. The cycle graph C_n is Hamiltonian.

2. Any graph obtained from C_n by adding edges is Hamiltonian
3. The path graph P_n is *not* Hamiltonian.




Figure 2.3.2 The Icosian game (from [Puzzle Musuem](#)) and its solution (from [Wikipedia](#))

The term Hamiltonian comes from William Hamiltonian, who invented (a not very successful) board game he termed the "icosian game", which was about finding Hamiltonian cycles on the dodecahedron graph (and possibly its sub-graphs).

The main thing you'll need to be able to do with Hamiltonian graphs is decide whether a given graph is Hamiltonian or not. Although the definition of Hamiltonian graph is very similar to that of Eulerian graph, it turns out the two concepts behave very differently. While Euler's Theorem gave us a very easy criterion to check to see whether or not a graph Eulerian, there is no such criterion to see if a graph is Hamiltonian or not. It turns out that deciding whether or not a graph is Hamiltonian is NP-complete, meaning that if we could solve that problem efficiently, then you could solve a host of other difficult problems efficiently as well.

It may seem unfair, then, to ask whether a graph is Hamiltonian or not. But it's only in a very theoretical way that the problem is extremely difficult -- as the number of vertices get very large, the problem gets harder and harder-quickly. For any given graph with a low number of vertices, there aren't that many possibilities to check.

If a graph *is* Hamiltonian, then by far the best way to show it is to exhibit a Hamiltonian cycle, as in [Figure 2.3.2](#). When the graph isn't Hamiltonian, things become more interesting.

The most natural way to prove a graph isn't Hamiltonian is to do a case by case analysis of possible paths, showing it doesn't work. For instance, in lecture we outlined the proof that if you remove a vertex from the Icosian graph, then the result isn't Hamiltonian. A natural way to do this is to pick a vertex, and consider the possible pairs of edges the path might take through that vertex. For each possibility, we know some edges won't be used, and can go further along that way.

In general, brute-force case-by-case analyses are proofs we want to avoid when possible, because it can be difficult to make sure we have actually found all the cases, and the proofs aren't always enlightening. It's much better when we can find a "reason" why the graph isn't Hamiltonian.

Example 2.3.3




Figure 2.3.4 A local configuration that can't exist in a Hamiltonian graph

Figure 2.3.4 shows a portion of a larger graph \mathbf{G} . The exact number of other vertices in the graph that X and Y are adjacent to is not important; what matters is that A and B are each adjacent to two vertices, X and Y . Any path through A would have to use X and Y , but so would any path through B . But then we have a small four cycle $XAYBX$ which doesn't use any other vertices in the graph, and so \mathbf{G} cannot be Hamiltonian. \square

Lemma 2.3.5 *Suppose that \mathbf{G} is bipartite and Hamiltonian, with n red vertices and m blue vertices.*

Proof. Consider a Hamiltonian path in \mathbf{G} . Since every edge is between a red and blue vertex, the vertices in the path must alternate between red and blue. Considering every other edge of the cycle pairs each red vertex with a blue vertex, and hence $n = m$. \blacksquare

The contrapositive of Lemma 2.3.5 can be used to show graphs *aren't* Hamiltonian: if \mathbf{G} is bipartite but doesn't have the same number of red vertices and blue vertices, then it can't be Hamiltonian.

Lemma 2.3.6 *The Petersen Graph is not Hamiltonian*

Proof. Of course, a case-by-case analysis of possible Hamiltonian cycles is possible. The number of cases can be reduced by using symmetries of the Petersen graph. Instead, for variation and to illustrate a different proof technique, we will use another method.

Assume for contradiction that the Petersen graph is Hamiltonian, and draw the ten vertices v_1, v_2, \dots, v_{10} around the cycle. The Hamiltonian cycle uses 10 of the 15 edges in the Petersen graph, and so there must be 5 more edges, with each vertex incident to one of them, as in the Petersen graph every vertex has degree 3.

Let's analyse where else the edge adjacent to v_1 could go. It can't go to v_1 itself, as the Petersen graph has no loops, and it can't go to v_2 or v_{10} as the Petersen graph has no multiple edges. If it went to v_3 or v_9 it would make a three cycle, which the Petersen graph doesn't have, and if it went to v_4 or v_7 , there'd be a four cycle. Hence, the only options are v_1 is adjacent to v_5, v_6 or v_7 . By reversing the direction of the Hamiltonian path, v_5 and v_7 are equivalent, and there are only two cases.

But the same analysis holds for every vertex: the extra edge at any vertex can either go to the opposite side of the circle, or be "off by one" and skip three vertices to either direction.

We now claim that not all the extra edges can go "directly across" -- there must be at least one edge that's off by one. If all the extra edges went directly across, then v_1 would go to v_6 , and v_2 would go to v_7 , and $v_1 - v_6 - v_7 - v_2 - v_1$ would be a 4 cycle.

Hence, without loss of generality we may assume that the extra edge at v_1 is $v_1 - v_5$. Let us then consider the extra edge at v_6 . It can't go directly across, as that is v_1 which already has its extra edge. Hence it must be off

by one, and go to either v_2 or v_{10} . But either way we get a four cycle: either $v_1 - v_2 - v_6 - v_5 - v_1$, or $v_1 - v_{10} - v_6 - v_5 - v_1$. Hence, we have a contradiction, and the Petersen graph cannot be Hamiltonian. ■

Finally, Ore's Theorem, a positive result, giving conditions which guarantee that a graph has a Hamiltonian cycle. First, a little bit of intuition. If we take an edge to a Hamiltonian graph the result is still Hamiltonian, and the complete graphs K_n are Hamiltonian. Thus, one might expect that a graph with "enough" edges is Hamiltonian.

The trick is in finding an interesting meaning of the word "enough". Simply counting the number of edges does not give very interesting bounds on what "enough" means, however -- the complete graph has $n(n-1)/2$ edges, but we can make it non-Hamiltonian by removing only $n-2$ edges: simply pick a vertex v and remove all but one of the $n-1$ edges coming out of v ; then v will now have degree 1, and hence the resulting graph cannot be Hamiltonian.

Theorem 2.3.7 Ore's Theorem. *Let G be a simple graph with n vertices, and assume that whenever two distinct vertices v, w are not adjacent, we have $d(v) + d(w) \geq n$. Then G is Hamiltonian.*

Proof. We will argue by contradiction, and begin by passing to a maximal counterexample. Note that if G satisfies the hypotheses, and we add an edge e to G between two non-adjacent vertices v and w , then the result will still satisfy the hypothesis. Indeed, we've only increased the degree of some vertices. So, we had a counterexample G to Ore's Theorem, we could iteratively add edges to G that didn't create Hamiltonian cycles, until we got a graph G that satisfies the hypotheses of Ore's theorem, doesn't have any Hamiltonian cycles, but if we add any edge e to G the result is Hamiltonian.

We now observe that such a G must have a Hamiltonian *path*: indeed, pick any edge e not in G and add it to G . The resulting graph is by assumption Hamiltonian, and since G wasn't Hamiltonian, the Hamiltonian cycle must contain the edge e . Deleting the edge e from the Hamiltonian cycle gives a Hamiltonian path in G .

Thus, let $v_1 - v_2 - \dots - v_n$ be a Hamiltonian path in G . We know v_1 and v_n are not adjacent, as otherwise G would be Hamiltonian. Thus, since G satisfies the hypothesis of Ore's theorem, we know $d(v_1) + d(v_n) \geq n$. We already have one edge adjacent to v_1 , and one edge adjacent to v_n , and so there must be at least $n-2$ other edges adjacent to one or other of these vertices. We will see that no matter how we add $n-2$ edges to these two vertices, we will create a Hamiltonian cycle.

To see this, note there is ever an i with v_1 adjacent to v_i and v_{i-1} adjacent to v_n , then G would have a Hamiltonian cycle: namely $v_1 - v_2 - \dots - v_{i-1} - v_n - v_{n-1} - \dots - v_i - v_1$. Now, there are $n-3$ different vertices we can add edges to v_1 to, namely $v_3 - v_{n-1}$, and similarly there are $n-3$ vertices we can add edges connecting v_n to, namely v_2, \dots, v_{n-2} . We arrange these $2(n-3)$ edges into a grid with 2 rows and $n-3$ columns, so that the two edges in each column are $v_1 - v_i$ and $v_n - v_{i-1}$, a pair of edges that can form a Hamiltonian cycle as in the last paragraph.

As we need to add at least $n-2$ edges, but we only have $n-3$ columns, there must be at least one column that contains two edges by the pigeonhole principle, but then we can create a Hamiltonian cycle using those two edges and the edges in our Hamiltonian path. ■

Note that Ore's Theorem is *not* an if and only if, and so Ore's Theorem cannot be used to prove that graphs *aren't* Hamiltonian. Indeed, there are plenty of graphs that are Hamiltonian but do not satisfy the hypotheses of Ore's Theorem. For instance, the cycle graph C_n is Hamiltonian, but every

vertex has degree 2, so if $n \geq 5$ the hypotheses of Ore's Theorem are not satisfied.

We also highlight that the proof began by considering a maximal counterexample to Ore's Theorem, and considering maximal/minimal counterexamples is often a useful proof technique, as you the maximality/minimality gives you some extra structure to work with.

Exercises

1. The questions in this exercise pertain to the graph G shown in [Figure 2.4.1](#).

- What is the degree of vertex 8?
- What is the degree of vertex 10?
- How many vertices of degree 2 are there in G ? List them.
- Find a cycle of length 8 in G .
- What is the length of a shortest path from 3 to 4?
- What is the length of a shortest path from 8 to 7?
- Find a path of length 5 from vertex 4 to vertex 6.




Figure 2.4.1 A graph

- Draw a graph with 8 vertices, all of odd degree, that does not contain a path of length 3 or explain why such a graph does not exist.
- Find an eulerian circuit in the graph G in [Figure 2.4.2](#) or explain why one does not exist.




Figure 2.4.2 A graph G

- Consider the graph G in [Figure 2.4.3](#). Determine if the graph is eulerian. If it is, find an eulerian circuit. If it is not, explain why it is not. Determine if the graph is hamiltonian. If it is, find a hamiltonian cycle. If it is not, explain why it is not.




Figure 2.4.3 A graph G

Chapter 3

Graphs on Surfaces

This chapter covers drawing graphs on surfaces. To motivate this topic, we will begin by thinking about videogames

We start with discussing whether or not graphs are planar, proving that $K_{3,3}$ and K_5 are not planar using a method we call the Planarity Algorithm for Hamiltonian graphs. We discuss the more general Kuratowski's theorem for proving any graph is planar or not. We introduce other surfaces, and how to draw graphs on them -- the sphere, Mobius band, and torus in particular. After a brief discussion of dual graphs, we prove Euler's theorem about planar graphs and explore several applications.

Introduction to Graphs on Surfaces

We begin our study of graphs on surfaces with an old chestnut of a problem, the solution of which we will develop into a more general algorithm.

The Utilities Problem

Suppose there are three houses, owned by Alice, Bob, and Carol, and they'd each like to be connected to one of three Utilities, say, gas, electricity, and water. There is no real difficulty in the real world, but if we add the restriction that we don't want any of the lines to cross over or under each other, the problem becomes quite interesting. A failed attempt at drawing a solution is shown here.




Figure 3.1.1 An attempt at solving the three utilities problem

Although this attempt failed, it seems very difficult to rule out that some other attempt wouldn't succeed; trying to make a case by case argument seems quite difficult to organize, and it's not clear that there are even finitely many possibilities. We need a careful way to approach the problem, which we will do in a moment, but first we will use this problem as motivation to make a few definitions.

Definition 3.1.2 A graph is *planar* if it can be drawn on a piece of paper so that no edges cross. \diamond

That definition is a bit loose -- for instance, it's left implicit, we're drawing the edges as lines, with the endpoints being the two vertices it connects. But this will be strong enough for our purposes.

With this definition in hand, the Utilities Question is asking whether the graph $K_{3,3}$ is planar -- treat the three utilities as red vertices, say, and the three houses as the blue vertices. This doesn't really help us organize our proof, however. To do that, we will use the basic fact that any circle drawn on the plane has an inside and an outside.

This last fact sounds absolutely trivial, but first, it is not true on other surfaces, for instance, on the torus -- in our video game world, the top of the screen makes a circle, but a point just above this circle is really at the bottom of the video game world, and so the circle doesn't cut the torus into two pieces; I also illustrated this with the Mobius band: the central line down the middle doesn't separate it into two pieces. This fact is usually stated as follows:

Theorem 3.1.3 Jordan Curve Theorem. *Any simple closed plane curve has an interior and an exterior*

Though easy to state, and intuitively obvious, the Jordan Curve Theorem is surprisingly subtle and difficult to prove; we won't use any more topology than this.

Before seeing it in practice, let's discuss informally how the Jordan Curve Theorem can be used to help prove whether a graph G is planar or not. Suppose that we have found a large cycle C_k as a subgraph of G . Then, if we had a planar drawing of G , this cycle would have to appear as a circle. By the Jordan Curve Theorem, this circle would have an inside and an outside, and every vertex and edge not in our cycle C_k would have to be either entirely within the circle, or entirely outside the circle. This gives us a way to organize the case by case argument.

The bigger a cycle we can find, the fewer other vertices and edges we need to consider, and so we have a much cleaner case by case argument. In the best cases, the graph is Hamiltonian and the cycle C_k includes all the vertices of G , and we only have to do a case by case analysis for some the remaining edges.

Let's see how this general principle gets illustrated in practice

Theorem 3.1.4 $K_{3,3}$ isn't planar

Proof. First let's name the vertices of $K_{3,3}$: let the vertices a, b, c be the blue circle vertices, and x, y, z be the red rectangle vertices. Then the path $a - x - b - y - c - z - a$ is a Hamiltonian cycle, and so if $K_{3,3}$ were planar it would be drawn as a circle in the plane, as shown below:




Figure 3.1.5 The Hamiltonian cycle in $K_{3,3}$

This contains 6 of the 9 edges of $K_{3,3}$; we need to add the edges $a - y$, $b - z$ and $c - x$. The edge $a - y$ could be drawn inside the circle or outside, suppose we draw it inside, as shown below, with the added edge dashed.




Figure 3.1.6 Adding $a - y$ inside

Then on the inside of the circle, x and c are on different sides of the line $a - y$, and so the edge connecting them must go outside the circle. The added edge could go around the right of the circle, as shown below here:




Figure 3.1.7 Adding $a - y$ inside
or around the left, as shown here:




Figure 3.1.8 Adding $a - y$ inside

But now b and z are different sides of $a - y$ inside the circle, and on different sides of $c - x$ outside the circle, and so cannot be connected without making two edges cross.

If we had begun by drawing $a - y$ outside the circle, then we would have had to draw $c - x$ inside the circle, and had the same problem with being able to draw the last line; as shown here:




Figure 3.1.9 Adding $a - y$ inside

■

The planarity algorithm for Hamiltonian graphs

In the previous chapter we showed that $K_{3,3}$ isn't planar; in this section we investigate how the ideas we used to solve the utilities problem for $K_{3,3}$ -- namely, the Jordan Curve theorem and the fact that $K_{3,3}$ is Hamiltonian -- generalize to other graphs. In the end, this will culminate in "The Planarity Algorithm for Hamiltonian Graphs".

Stereographic Projection and Unnecessary cases

It will make our life easier if before we investigate other graphs we streamline our proof for $K_{3,3}$ slightly: there were a few times where we had to treat different cases that wound up behaving essentially the same, and we'd like to see that we didn't actually need to treat them as separate cases. In particular, we would like to show that the following three seemingly different ways to connect the first two vertices lead to the same analysis:

1. Connecting them inside the Hamiltonian cycle
2. Connecting them outside "to the left"
3. Connecting them outside "to the right"




Figure 3.2.1 The three cases

The solution will be to think about drawing the graphs on the sphere S^2 instead of the plane. First, let's see why this solves our problem. On the plane, the inside of a circle is different from the outside of a circle -- the inside is bounded, but the outside is unbounded. However, on the sphere, the two sides of a circle are equivalent -- you can deform any circle to be an equator, and then the northern hemisphere is equivalent to the southern hemisphere. This shows on there sphere, the inside and the outside aren't really different cases.

Furthermore, going around the outside to "the left" or "to the right" are equivalent on the sphere -- you can slowly make the path around the sphere bigger and bigger, and then slip it around the north or south pole, and back. Alternatively, we've already seen that the inside of the circle is equivalent to the outside of the circle on the sphere S^2 , and on the inside of the circle it doesn't matter exactly how the two points are connected, and so it shouldn't matter on the outside, either.

So we've argued that if we're trying to draw a graph on the sphere, all three cases are the same, but it should still feel like a bait-and-switch: we weren't trying to draw graphs on there sphere, we were trying to draw graphs on the plane. The connection comes from the fact that the sphere can be viewed as a plane with one additional point.

Proposition 3.2.2 *Let $p \in S^2$ be any point. Then $S^2 \setminus \{p\} \cong \mathbb{R}^2$.*

Proof. One way to visualize this is imagine the sphere as being made from very flexible clay. If we poke a small hole in the top of the sphere, we could stick our fingers in and make the hole larger, and gradually stretch and bend and reform for the sphere to be a flat disk, which could be stretched to be the whole plane, in the same way the tangent function maps the interval $(-\pi/2, \pi/2)$ to the whole real line \mathbb{R} .

Alternatively, one could use stereographic projection, as shown in Figure . Draw S^2 in \mathbb{R}^3 as the unit sphere at the origin, and let $N = (0, 0, 1)$ be the north pole of the sphere. Stereographic projection gives a bijection between $S^2 \setminus \{N\}$ (the sphere minus the north pole) to the plane, as follows: for any point $p \neq N$ the line through p and N must meet the xy -plane at one point. On the other hand, any line through N and a point on the xy -plane must meet the sphere at one other point.




Figure 3.2.3 Stereographic Projection

Accepting that S^2 is \mathbb{R}^2 minus one point, we see that we can draw a given graph G on S^2 if and only if we can draw G on \mathbb{R}^2 : if we draw it on \mathbb{R}^2 , we can view the \mathbb{R}^2 as a small patch of S^2 . And if we have a drawing on S^2 , there must be at least one point on S^2 that isn't in the drawing of G , and doing

stereographic projection from that point gives a drawing of G on the plane \mathbb{R}^2 . ■

The planarity algorithm

The preceeding discussion may have felt heavy going, but the upshot is that the cases that seemed "the same" in our analysis of $K_{3,3}$ actually are the same, and similar cases will be the same for any graph. This will make it much easier to extend our reasoning to more complicated graphs.

Suppose that G is Hamiltonian, and choose a Hamiltonian cycle; if G were planar then this cycle must be drawn as a circle, and every other edge must either lie entirely inside or entirely outside the graph. Now consider two edges $e = ab$ and $f = xy$ that are not part of the cycle. Depending on the order that a, b, x and y appear as you go around the Hamiltonian cycle, one of two things will happen:

1. If the vertices of e and f do not interlace (e.g. $abxy, yxab, xbay, \dots$), or if they share a vertex (e.g., $a = x$), then e may be drawn both inside or both outside the circle without crossing
2. If the vertices of e and f do interlace (e.g. $axby, xayb, yaxb, \dots$) then if e and f are drawn both inside or both outside the circle, they must cross

This motivates the following definition

Definition 3.2.4 $\text{Cross}(G, C)$. Let G be a Hamiltonian graph, and C a Hamiltonian cycle in G . The *crossing graph* of G and C , denoted $\text{Cross}(G, C)$ has as vertices the edges of G that aren't in the cycle, and an edge between vertices p and q if the vertices of the corresponding edges interleave -- that is, p and q are adjacent if they cannot be drawn on the same side of the cycle C without crossing. ◇

Algorithm 3.2.5 *The planarity algorithm for complete graphs.* Suppose that G is Hamiltonian, and C is a Hamiltonian cycle. Then G is planar if and only if $\text{Cross}(G, C)$ is bipartite.

The idea is that if G is planar, the vertices of $\text{Cross}(G, C)$ are naturally bicolored, with the red vertices, say, corresponding to the edges that are drawn inside the cycle C , and the blue edge corresponding to the edges that are drawn outside the cycle C . The definition of the edges of $\text{Cross}(G, C)$ guarantees there are no edges between vertices of the same color.

Similarly, if we can find a colouring of the vertices of $\text{Cross}(G, C)$ so that adjacent vertices have different colours, then we can draw all the edges of G corresponding to red vertices of $\text{Cross}(G, C)$ inside (or outside) C without having any of them cross.

Example 3.2.6 *The complete graph K_5 isn't planar.* Let's apply the planarity algorithm to the complete graph K_5 . Let's label the vertices 1-5, and take our Hamiltonian cycle C to be 123451, which we've drawn as the outside pentagon in the following figure:




Figure 3.2.7 The graph (K_5)

Since K_5 has $\binom{5}{2} = 10$ edges, there are 5 edges that aren't used in C , namely 13, 14, 24, 25, and 35. So $\text{Cross}(K_5, 123451)$ will consist of these five vertices. We see that 13 will be adjacent to 24 and 25, since these edges would cross if drawn inside, but 13 is not adjacent to 14 or 35, since these edges would cross 13 if drawn on the same side of the circle, as illustrated in the next figure




Figure 3.2.8 The edges 13 does and does not cross

Similar consideration with the other edges show that $\text{Cross}(K_5, 123451)$ is the following graph, which is isomorphic to a five cycle:




Figure 3.2.9 The graph $\text{Cross}(K_5, 123451)$

In particular, $\text{Cross}(K_5, 123451)$ is not bipartite. Hence, by the planarity algorithm for Hamiltonian graphs, we see that K_5 is not planar. \square

Example 3.2.10 A planar graph. Let's use the planarity algorithm for Hamiltonian graphs to find a planar drawing of the graph shown in the next figure.




Figure 3.2.11 A graph H

We see that H is Hamiltonian and take as our Hamiltonian cycle the path around the outside, namely $abcxyza$. There are then six edges not contained in the Hamiltonian cycle, and we find that $\text{Cross}(H, abcxyza)$ is as follows:




Figure 3.2.12 The graph $\text{Cross}(H, abcxyza)$

For instance, in H the edge ax crosses the three edges cy, cz and bz , and so in $\text{Cross}(H, abcxzya)$, the vertex ax is adjacent to these vertices.

The graph $\text{Cross}(H, abcxzya)$ has no odd cycles and hence is bipartite -- for instance, we may color ax, bx and ay red, and the other three vertices blue. Then, to draw H in the plane without edges crossing, we draw the red edges inside the cycle, and the blue edges outside the cycle:




Figure 3.2.13 The graph H drawn without edges crossing

□

Kuratowski's Theorem

The planarity Algorithm for Hamiltonian graphs gives a very convenient and systematic way to determine whether a Hamiltonian graph is planar or not, and we saw that with some work it can be hacked to work for graphs that are "almost" Hamiltonian -- that have a cycle that go through all but one or two vertices, say.

Stretching these ideas further, the general logic of considering cycles and applying the Jordan Curve theorem to them would provide a way to prove whether an arbitrary graph is planar or not, but as we have more or more vertices that aren't on our cycle to consider the arguments would get more and more complicated, and it's clear that a better method is desirable. In this section we will present, (but not completely prove) Kuratowski's theorem, which gives a method to determine whether or not an arbitrary graph is planar.

Planarity, subgraphs, and subdivisions

The idea behind Kuratowski's theorem rests on two small observations, which we illustrate in a simple example before discussing more formally.




Figure 3.3.1 Two nonplanar graphs

The graphs G_1 and G_2 in Figure 3.3.1 both *look* a lot like $K_{3,3}$, and since $K_{3,3}$ is nonplanar, we might expect them to be nonplanar as well, but we need to be careful and precise in checking this. We work this out in the next example.

Example 3.3.2 Two nonplanar graphs. The graph G_1 in Figure 3.3.1 is $K_{3,3}$ with a vertex and two edges added to it; put another way, $K_{3,3}$ is a subgraph of G_1 . If we could draw G_1 in the plane, we could just ignore this extra vertex and these two edges, and we'd have a drawing of $K_{3,3}$ in the plane, but we know $K_{3,3}$ isn't planar. So to avoid contradictions we are forced to conclude that G_1 isn't planar.

The graph G_2 looks just like $K_{3,3}$, but we have added an extra vertex of degree two in the middle of one of the edges. Note that $K_{3,3}$ is *not* a subgraph of G_2 , and so we need to use slightly different reasoning than we did for showing G isn't planar. But drawing G_2 is just like drawing $K_{3,3}$ and then adding an extra dot for the new vertex of degree two. If we could draw G_2 in the plane, we could just skip adding this extra dot, and we'd have a drawing of $K_{3,3}$ in the plane. Again, since we know $K_{3,3}$ isn't planar, we see that G_2 isn't planar, either. \square

We will now generalize the methods we used to show G and H are nonplanar and summarize them as lemmas. The reasoning we used to prove that G was nonplanar doesn't need to be changed at all to prove:

Lemma 3.3.3 *If H is nonplanar, and H is a subgraph of G , then G isn't planar.*

Proof. We've essentially already proved it, but we'll restate the reasoning in a different way for completeness.

When we draw a graph G in the plane, we also draw all the subgraphs of G in the plane. Thus, if G is planar, then all of its subgraphs are planar. Our lemma is the contrapositive of this statement. \blacksquare

Lemma 3.3.3 is logically equivalent to the discussion above, but it's worth restating the logic in this direction as well: if we can't draw G in the plane, then we certainly can't draw H in the plane without edges crossing, as if we could then we'd have a drawing of H as part of our big drawing.

Example 3.3.4 Complete graphs. If $n > 5$, then K_n is not planar by Lemma 3.3.3, because K_5 is a subgraph of K_n , and we know that K_5 isn't planar.

We could also have used the fact that $K_{3,3}$ is a subgraph of K_n , and $K_{3,3}$ is also nonplanar. \square

To generalize the method we used to prove G_2 is non-planar, we first make a formal definition that encapsulates the idea of "adding dots" to the middle of edges:

Definition 3.3.5 Subdivision. We say that G is a *subdivision* of H if G is obtained from H by adding some vertices of degree two in the middle of some of the edges of H . \diamond

Example 3.3.6 Cycles and Paths. Any path graph $P_n, n \geq 2$ is a subdivision of the graph P_2 consisting of two vertices with an edge between them. Any cycle graph $C_m, m \geq 3$ is a subdivision of the triangle C_3 . \square

Lemma 3.3.7 *If G is a subdivision of a nonplanar graph H , then G is nonplanar.*

Proof. Suppose that G was planar, and draw it in the plane. Then erase the vertices of degree two we added when we subdivided H , merging the edges on either side to one. We obtain a planar drawing of H , a contradiction, and so G must not have been planar. \blacksquare

Kuratowski's Theorem

The definitions and lemma of the previous section essentially prove the "easy" direction of the following theorem, which will be our main tool for proving graphs aren't planar.

Theorem 3.3.8 Kuratowski's Theorem. *A graph G is nonplanar if and only if G has a subgraph that's a subdivision of $K_{3,3}$ or K_5 .*

Proof. We will only prove one direction: that if G has such a subgraph, then G is nonplanar; the other direction is more difficult.

Suppose that H is a subgraph of G that is a subdivision of $K_{3,3}$ or K_5 . Since we've proven $K_{3,3}$ and K_5 are nonplanar, we know H is nonplanar by [Lemma 3.3.7](#). Now since H is a subgraph of G and we know H is nonplanar, we know G is nonplanar by [Lemma 3.3.3](#). \blacksquare

Although we've only proven one direction of Kuratowski's theorem, it's the important direction -- the one that lets us prove graphs are nonplanar. The other direction would tell us information about subgraphs of a graph that we already knew was nonplanar for some other reason. Or, taking the contrapositive, it would let us prove a graph *was* planar by looking at all subgraphs of it and showing none of them looked like K_5 or $K_{3,3}$. But this would be a lot of work and there's a much easier way to show a graph is planar: draw it in the plane! If you're asked to prove a graph is planar, you will almost always also be asked to draw it in the plane.

However, we will *implicitly* use the hard direction in the following way: if a graph G is nonplanar, you can *always* use Kuratowski's theorem to prove that it's nonplanar. This is reassuring because it means our tool will always work to prove it's nonplanar, and that you aren't wasting your time looking for subgraphs that don't exist.

Applying Kuratowski's Theorem

The tricky part of using Kuratowski's theorem is actually *finding* the desired subgraph. We won't really discuss algorithm aspects of this; for any graph you are asked to prove non-planar, it will be possible to do so by educated trial and error. A few rules of thumb may be helpful, however. First, note

that subdivision cannot increase the degree of any vertex. So, for \mathbf{G} to have a subgraph that's a subdivision of K_5 , \mathbf{G} has to have at least 5 vertices with degree at least 4; if it doesn't, but we still suspect \mathbf{G} to be nonplanar, we know instead that we should be looking for a subdivision of $K_{3,3}$.

Conceptually, it can be useful to think that some vertices of \mathbf{G} are going to be vertices of your K_5 or $K_{3,3}$, and we're going to need to connect those. We can use the remaining vertices of \mathbf{G} as parts of subdivided edges between these, but these extra vertices can only be used in at most one such connection. Thus, these extra vertices are a limited resource we have, and a useful heuristic in looking for subgraph is to take a "greedy" approach, where we choose our vertices to require as few subdivisions as possible to make connections. We illustrate this idea in the next example.

Example 3.3.9 The Petersen graph isn't planar.




Figure 3.3.10 The Petersen graph, labeled

Let us use Kuratowski's Theorem to prove that the Petersen graph isn't planar; Figure 3.3.10 has a drawing of the Petersen graph with the vertices labeled for reference. Since the Petersen graph is regular of degree three, we know that it can't have a subgraph that's a subdivision of K_5 , as it would need to have some vertices of degree 4 or larger.

It makes sense to attempt a greedy algorithm -- in the standard drawing of the Petersen graph, pick the very top vertex 1 to be "red" and the three vertices adjacent to it to, 2, 5, and 6, to be "blue". We need two more red vertices. All vertices left are adjacent to exactly one blue vertex, so from a greedy point of view there is no preference for which vertex we pick to be the next blue vertex. Let us pick 9 to be another red vertex. Then it is connected directly to blue vertex 6, but we must find paths from 9 to 2 and 5. We could, for instance, take the path 9-4-3-2 to connect to 2, but that would use two vertices up while

the path 9-7-2 only uses one extra vertex, and so seems better. Then we can connect 9 to 5 through vertex 4, and vertex 9 has been connected to all the blue vertices.

Now, we need to choose one more vertex to be a red vertex, and the vertices we haven't used are 3, 10, and 8. If we tried to make 3 the last red vertex we run into a problem: we need to connect vertex 3 to 3 other vertices, but one of the edges goes to vertex 4 which was one of the subdivided vertices that we can't visit again. Hence, we only have two possible paths out of 3, and will ever only be able to connect it to two blue vertices. A similar problem occurs if we try to make 10 the last red vertex -- it's adjacent to the vertex 7 used as a subdivided vertex. The remaining choice is vertex 8, which works, as shown in the following diagram.




Figure 3.3.11 A subdivided $K_{3,3}$ in the Petersen graph

The red and blue vertices of the subdivided $K_{3,3}$ are shown as squares/circles, and the edges of the subdivided $K_{3,3}$ are colored thick -- only the dotted edges 7-10 and 3-4 of the Petersen graph are not used in the subgraph.

□

Drawing Graphs on Other surfaces

We saw, using stereographic projection, that being able to draw a graph on the sphere is the same as being able to draw the graph on the plane. In this section we will discuss drawing graphs on other surfaces -- the torus and the Möbius strip we will discuss in detail, though similar ideas work for any surface. We need a way to represent such graphs on a piece of paper, for use in a book (or on the exam, say). Much of the material from the rest of this chapter (Kuratowski's theorem, Euler's theorem) have analogues for other surfaces, but are beyond the scope of this module.

Motivation and culture: Manifolds and Surfaces

In this short subsection we are going to be slightly informal. The goal is simply to motivate this section about drawing graphs on surfaces other than the sphere, and to give a motivating problem that we will solve the next section on Euler's Theorem.

People used to think that the earth was flat, because if you can't see the whole thing, but can just look at just one little patch of it around you, it looks like a piece of \mathbb{R}^2 . Formally, a mathematicians say "locally homeomorphic to \mathbb{R}^2 " to mean that you can't tell it apart from \mathbb{R}^2 by just looking at a small piece of it.

To compliment the familiar earth idea of the earth being round, we give a few more shocking thought-experiment examples. The first is: how do we know the earth is round? I, personally haven't been to space, haven't been all over the world. Maybe there's a giant tunnel running from the south pole to the north pole, and the earth is really a *torus* (the surface of a donut or a bagel).

The idea of the earth being the surface of the torus probably seems absolutely ridiculous, but it happens "by accident" in videogames. In old games like Pacman or Asteroids, the game takes place on one computer screen, but to keep it from having edges and corners the designers made the game "wrap around" -- if anything goes off the right edge of the screen, it comes back at the corresponding spot on the left edge of the screen, and similarly anything that goes off the top of the screen comes back on the bottom of the screen.

A similar model, expanded slightly, is used in many other video games, like early ones in the Final Fantasy series, use essentially the same process to model the surface of a planet. We claim that any of these videogame worlds are actually the surface of a torus. Instead of a computer screen, let's put the world on a very flexible flat sheet. To get the left and right edges to match up, we can curl the screen up into a cylinder -- going off one edge takes us around the back of the cylinder. Now, our cylinder has two circular boundaries, coming from the top and bottom of the screen, and to get these to match up we can bend our cylinder up and glue these together to make a torus.

Finally, we can step up a dimension, we generally assume that the three dimensional space we live in is \mathbb{R}^3 , but what evidence do we have for that? Maybe it's got some different shape, and if we could fly for untold light years in one direction we'd come back to the earth from a different direction! I can recommend Janna Levin's popular science book / memoir "How the Universe got its Spots" for an account of how physicists studied whether patterns in the cosmic microwave background radiation could have been created by the universe being a shape other than \mathbb{R}^3 .

That discussion may feel out of place; its purpose was to motivate the following definition, which we will then apply to graph theory:

Definition 3.4.1 An *n-dimensional manifold* is a space that is locally homeomorphic to \mathbb{R}^2 . A *surface* is a two dimensional manifold. \diamond

Graphs on the Torus

The *torus* is another word for the surface of a donut. There are some graphs that can't be drawn on the sphere, but can be drawn on a torus. But we need a way to represent drawing graphs on the torus just using a normal sheet of paper -- it would be awkward and impractical to hand every student an innertube or a donut at the exam to hand in with their papers.

We will do this by copying the videogame worlds we saw in the introduction. Draw a square to represent the videogame screen, and then draw the graph inside the square, with the added proviso that if while drawing an edge of the graph we hit the border of the square, we continue the edge at the same point of the opposite side of the square. To make clear what we're doing, it is useful to draw arrows on the edges of the square to indicate how they are being identified -- we put one arrowhead, pointing the same direction, on the left and right edges, to indicate that they are being identified, with the tip end of one edge being marked to the tip end of the other, and the tail end being matched to the tail end of the other. We draw two arrowheads on the top and bottom edges, also pointing the same direction.




Figure 3.4.2 K_5 and K_6 drawn on a torus

This whole process is illustrated in [Figure 3.4.2](#), where the complete graphs on five and six vertices are drawn on the torus. We end with a few short tips and tricks for trying to draw graphs on surfaces.

The first is that it can be complicated to keep track of edges that wrap around the sides of the torus. If lots of them are wrapping around, it can be easy to lose track of which edges are connecting to which others on the opposite side. One way to make this clearer is to label the crossing point on one side letters or numbers in order, and then on the opposite side label with the letters and numbers going in the same direction as indicated by the arrow. Then you know that the edge that crosses at letter c , say, on one side, picks up at c on the other edge.

Another idea is to try to minimize the number of such crossings, and draw as much of the graph as possible in the center of the square, and only a few edges wrapping around the torus. One heuristic to follow is to work like we did with the planarity algorithm for Hamiltonian graphs, and try to draw a large cycle as a cycle in the square. Then, connect as many edges and extra vertices as possible through the centre of the cycle, and only when you've ruled that out try to wrap edges around the torus. This is the approach taken in the drawing of K_5 in [Figure 3.4.2](#); the outside house is a Hamiltonian cycle, two edges could be drawn in the center of the house, and then the remaining three edges have no choice to be drawn wrapping around the torus.

The drawing of K_6 using another trick that is worth explaining. It begins with the drawing of K_5 in black, and adds new stuff in red. It appears that four new vertices have been added, one at each corner of the square. But actually,

when the edges are identified, the four corners of the sphere get identified into one point on the torus. (Check this by visualizing what happens if you folded the paper up!). Placing one of the vertices on this graph on the corner can be useful because then edges drawn from this vertex will not need to wrap around the sides of the square.

Möbius strip

If you take a piece of paper and roll it up, identifying the two edges in the usual way, one obtains the cylinder. But if instead you identify the edges in the opposite way as usual, you one obtains a different surface called the *Möbius strip*. The Möbius strip is famous for only having one side. In the drawing shown below, if one of the ants walks all the way around the strip, when it returns to where it starts it will be on the opposite side of the strip. Similarly, whereas a cylinder has two boundary circles, the Möbius strip only has one.




Figure 3.4.3 M.C. Escher's Möbius Strip II

The fact that the Möbius strip only has one boundary circle has the following surprising consequence, that makes a great "party trick", is to make a Möbius strip by taping up a strip of paper, and then cut it down the very middle of the strip -- you wind up not getting two pieces of paper, but just one!

To actually represent graphs drawn on the Möbius strip, we work similarly to what we did for the torus; we draw a square, and then we draw arrows on the left and right edges to indicate that these edges are drawn together. However, we have one arrow drawn up, and one arrow drawn down, to indicate that the ends of the strip are glued together with a half twist. Since the top and bottom of the Möbius strip do not get identified together at all, we do not draw any arrows on them. Then, if an edge goes off one end the Möbius strip near the top, it comes back on the opposite end near the bottom, and vice versa.




Figure 3.4.4 Two drawings of $K_{3,3}$ on the Möbius strip

Finally, we explain some heuristics about drawing graphs on surfaces, with reference drawings of $K_{3,3}$ on the Möbius strip. As discussed at the end of the section of the torus, it can be useful to follow part of the reasoning used in the planarity algorithm for Hamiltonian cycles when trying to draw graphs on surfaces other than the sphere. That is, start with a large cycle in the graph -- for $K_{3,3}$, with a, b, c red vertices, and x, y, z the blue vertices, we will use the cycle $axbycza$. It may be that this cycle is drawn in the plane as a circle, as the in the left hand example; then we try to connect as many edges as possible through the centre of the circle, and then do the rest on the outside, possible wrapping around the Möbius strip.

But the Jordan Curve Theorem (that a circle has an inside and an outside) only holds for the sphere -- on other surfaces, there are always curves where this isn't true. Whether of necessity, or choice, it might be that our large cycle is drawn as one of these loops that doesn't have an inside or an outside. On the right hand side of Figure 3.4.4, we have drawn the cycle $axbycza$ as the centre of the Möbius strip.

Finally, we note that in general for a surface, there are multiple different ways to draw curves; in Figure 3.4.5, we have drawn our chosen Hamiltonian cycle as a curve that wraps *twice* about the Möbius strip. What happens if you cut a physical Möbius strip along this line?




Figure 3.4.5 A third drawing of $K_{3,3}$ on the Möbius strip

Euler's Theorem

This section covers Euler's theorem on planar graphs and its applications. After defining faces, we state Euler's Theorem by induction, and gave several applications of the theorem itself: more proofs that $K_{3,3}$ and K_5 aren't planar, that footballs have five pentagons, and a proof that our video game designers couldn't have made their map into a sphere without doing something very strange.

Counting faces

A **face** of a planar drawing of a graph is a region bounded by edges and vertices and not containing any other vertices or edges.

Figure ?? shows a planar drawing of a two graphs. The left graph has determines 5 regions, since we also count the unbounded region that surrounds the drawing.

*= *3.5.1)!

While it might seem like a coincidence that this computation results in 2 for these planar drawings, there's a more general principle at work here, and in fact it holds for *any* planar drawing of *any* planar graph.

In fact, the number 2 here actually results from a fundamental property of the plane, and there are a corresponding theorems for other surfaces. However, we only need the result as stated above.

Theorem 3.5.2 Euler's Formula. *Let \mathbf{G} be a connected planar graph with n vertices and m edges. Every planar drawing of \mathbf{G} has f faces, where f satisfies*

$$n - m + f = 2.$$

Proof. Our proof is by induction on the number m of edges. If $m = 0$, then since \mathbf{G} is connected, our graph has a single vertex, and so there is one face. Thus $n - m + f = 1 - 0 + 1 = 2$ as needed. Now suppose that we have proven Euler's formula for all graphs with less than m edges and let \mathbf{G} have m edges. Pick an edge e of \mathbf{G} . What happens if we form a new graph \mathbf{G}' by deleting e from \mathbf{G} ? If \mathbf{G}' is connected, our inductive hypothesis applies. Say that \mathbf{G}' has n' vertices, m' edges, and f' faces. Then by induction, these numbers satisfy

$$n' - m' + f' = 2.$$

Since we only deleted one edge, $n' = n$ and $m' = m - 1$. What did the removal of e do to the number of faces? In \mathbf{G}' there's a new face that was formerly two faces divided by e in \mathbf{G} . Thus, $f' = f - 1$. Substituting these into $n' - m' + f' = 2$, we have

$$n - (m - 1) + (f - 1) = 2 \iff n - m + f = 2.$$

Thus, if \mathbf{G}' is connected, we are done. If \mathbf{G}' is disconnected, however, we cannot apply the inductive assumption to \mathbf{G}' directly. Fortunately, since we removed only one edge, \mathbf{G}' has two components, which we can view as two connected graphs \mathbf{G}'_1 and \mathbf{G}'_2 . Each of these has fewer than m edges, so we may apply the inductive hypothesis to them. For $i = 1, 2$, let n'_i be the number of vertices of \mathbf{G}'_i , m'_i the number of edges of \mathbf{G}'_i , and f'_i the number of faces of \mathbf{G}'_i . Then by induction we have

$$n'_1 - m'_1 + f'_1 = 2 \quad \text{and} \quad n'_2 - m'_2 + f'_2 = 2.$$

Adding these together, we have

$$(n'_1 + n'_2) - (m'_1 + m'_2) + (f'_1 + f'_2) = 4.$$

But now $n = n'_1 + n'_2$, and $m'_1 + m'_2 = m - 1$, so the equality becomes

$$n - (m - 1) + (f'_1 + f'_2) = 4 \iff n - m + (f'_1 + f'_2) = 3.$$

The only thing we have yet to figure out is how $f'_1 + f'_2$ relates to f , and we have to hope that it will allow us to knock the 3 down to a 2. Every face of \mathbf{G}'_1 and \mathbf{G}'_2 is a face of \mathbf{G} , since the fact that removing e disconnects \mathbf{G} means that e must be part of the boundary of the unbounded face. Further, the unbounded face is counted twice in the sum $f'_1 + f'_2$, so $f = f'_1 + f'_2 - 1$. This gives exactly what we need to complete the proof. ■

Remark 3.5.3 Alternative method of dealing with the second case.

In our proof of Euler's theorem, the most complicated part was dealing with the situation if the edge e disconnects our graph \mathbf{G} when we remove it. In this case, instead of deleting the edge e we can contract it -- that is, shrink it to a point. This would have result in a graph that is still planar and still connected, but with one less edge (e is no longer around), and one less vertex (the two vertices e connects are now merged into one). The number of faces remains unchanged. So the number of edges and the number of faces each decreased by one, these two changes cancel out when we calculate $n - m + f$, and hence both are equal to two.

Applications of Euler's theorem

By itself, Euler's theorem doesn't seem that useful: there are three variables (the numbers of edges, vertices, and faces) and only one equation between them, so there are still lots of degrees of freedom. For it to be particularly useful, we want to have other relationships between these numbers. In many applications, these relationships can come from handshaking.

Recall that Euler's handshaking lemma said that

$$\sum_{v \in G} d(v) = 2|E(G)|,$$

the sum of the degrees of all the vertices is twice the number of edges. If we had some knowledge about the degrees of these vertices, we could get another relationship between the number of vertices and the number of edges. For example, if G is regular of degree k , then every vertex has degree k , and hence the sum of all the degrees is just kn . Hence, handshaking would tell us that $kn = 2m$, and we would have another relationship between the three variables m, n and f .

Similarly, there is a handshaking between faces and edges. Let the *degree* of a face be the number of edges that occur around it -- so, a triangle would have degree three. Then, if we sum up the degrees of all the faces, we're counting each edge twice again -- once from the face on its left, and once from the face on its right. so we have

$$\sum_{f \in \text{faces}(G)} d(f) = 2|E(G)|$$

Note that this is just the usual vertex-edge handshaking for the dual graph.

Thus, vertex-edge and face-edge handshaking can potentially give us two other sources of relationships between the numbers of vertices, edges, and faces. Most applications of Euler's theorem proceed by combining all three relationships, as we shall see.

Lemma 3.5.4 K_5 isn't planar

Proof. We give a proof by contradiction. Suppose K_5 was planar, and draw it on the plane. We know that K_5 has 5 vertices and $\binom{5}{2} = 10$ edges, and so by Euler theorem we know that any drawing of it must have 7 faces. We now use edge-face handshaking to get a contradiction.

What could the degrees of the faces be? We don't know for sure, but we know that none of the faces could have degree one or two, as then the edges would form a loop or multiple edges between two vertices, but K_5 is simple. Hence, every face must have $d(f) \geq 3$. But then handshaking gives:

$$20 = 2e = \sum d(f) \geq \sum 3 = 3 \cdot 7 = 21$$

which is the desired contradiction, and so we conclude that K_5 is not, in fact, planar. ■

It is a good exercise to adapt this proof to prove that $K_{3,3}$ isn't planar; one needs to use the extra fact that $K_{3,3}$ doesn't have any three cycles (why not?)

We now prove that footballs have 12 pentagons. More precisely, use the shorthand *football graph* to mean any trivalent graph drawn on the plane so that every face is a pentagon or hexagon. Then we have:

Theorem 3.5.5 The football theorem. *Let G be a football graph drawn on the plane, with P pentagonal faces, and H hexagonal faces. Then $P = 12$.*

Proof. Let V, E, F denote the number of vertices, edges, and faces of \mathbf{G} . Since every face is a hexagon or pentagon, we have $E = P + H$, and substituting this into Euler's theorem gives:

$$V - E + P + H = 2 \quad (3.5.1)$$

Now we turn handshaking. Since \mathbf{G} is trivalent, every vertex has degree three, and so vertex-edge handshaking becomes:

$$2E = 3V \quad (3.5.2)$$

Finally, since there are P pentagonal faces and H hexagonal faces, face-edge handshaking becomes:

$$2E = 5P + 6H \quad (3.5.3)$$

Multiplying
by six gives:

$$6V - 6E + 6P + 6H = 12$$

Multiplying (3.5.2) by two gives $6V - 4E = 0$, which we can use to eliminate V , giving

$$-2E + 5P + 6H = 12$$

Finally, using (3.5.3) we can eliminate both E and H in one go, being left with $P = 12$ as desired. ■

Finally, we prove that given some sensible restraints, video game designers cannot make a world map that's a sphere. A videogame world locally looks like a square grid -- with every vertex and face having degree four.

Theorem 3.5.6 The videogame theorem. *It is impossible to draw for a graph to be drawn on the sphere so that every vertex and every face has degree 4.*

Proof. Since every vertex has degree 4, vertex-edge handshaking gives $2E = 4V$, and since every face has degree 4, face-edge handshaking gives $2E = 4F$. Thus, we see $V = F = E/2$, and plugging this in gives:

$$V - E + F = E/2 - E + E/2 = 0$$

which contradicts Euler's Theorem. Hence, such a graph on a sphere is not possible. ■

Exercises

1. Show that the Petersen graph has a cycle that uses all the vertices but one. Give another proof that the Petersen graph is not planar by modifying the planarity algorithm for Hamiltonian graphs to deal with this extra vertex.
2. Draw $K_{4,4}$ and K_7 on the torus.

3. Use the planarity algorithm to show that the given graph \mathbf{G} is planar, and draw a plane graph isomorphic to it. Explain how you might obtain a non-planar graph by adding one extra edge, and for your non-planar graph, find a subgraph that's a subdivision of K_5 or $K_{3,3}$.




Figure 3.6.1 A Hamiltonian graph \mathbf{G}

- 4.




Figure 3.6.2 A Hamiltonian graph \mathbf{H}

Let \mathbf{H} be the graph shown in Figure 3.6.2. Show that \mathbf{H} is not planar in two ways: using the planarity algorithm, and using Kuratowski's theorem. Draw \mathbf{H} on the Möbius strip and on a torus without the edges crossing.

5. A connected plane graph has faces of degrees 3 and 10 only, and every vertex has degree at least 3. Prove that it must have fewer faces of degree 10 than of degree 3. If every vertex has degree 3, prove that the number of faces of degree 10 must be a multiple of 3 and the number of faces of degree 3 must be a multiple of 4.

As an example, how many faces of degree 3 and degree 10 does the truncated dodecahedron possess? (A truncated dodecahedron is obtained by sliving off each vertex of a dodecahedron to give a triangle.)

Chapter 4

Colourings

This chapter covers several types of colouring questions on graphs. The initial motivation for these questions comes from an early question about colouring the countries on maps so that adjacent countries have different colours.

Chromatic number

The study of graph colourings began with the colouring of maps. Usually on a map, different regions (countries, counties, states, etc.) are visually distinguished from each other by giving each one a different colour, with the idea that adjacent regions should have different colours so that boundaries can be easily seen. For instance, in this old road map of England and Wales, each county is coloured either red, yellow blue or green, and bordering counties have different colours.



Figure 4.1.1 A historical example of a map colouring. [Image](#) courtesy Cartography Associates under a [creative commons license](#)

Note that in the map above, only four colours are used. In 1852 Francis Guthrie suggested that for any possible map drawn on a piece of a paper, four colours would be enough. Guthrie's conjecture wasn't proven for more than a hundred years later.

Let's make Guthrie's conjecture precise and connect it to graph theory. Note that as in Example , it makes sense to work with essential the dual

picture to the map; we make the regions of the map into vertices, and we put an edge between two regions if they share a vertex. Then we are lead to the following definitions.

Definition 4.1.2 Colourings and Chromatic number. Let \mathbf{G} be a graph. A k -colouring (or sometimes *vertex colouring*) of \mathbf{G} with k colours is an assignment of one of k colours to each of the vertices of \mathbf{G} so that adjacent vertices have different colours.

More formally, a k -colouring is a function $f : V(\mathbf{G}) \rightarrow \{1, \dots, k\}$ so that if $v \sim w$ then we have $f(v) \neq f(w)$.

The *chromatic number* of a graph \mathbf{G} , written $\chi(\mathbf{G})$, is the least number of colours needed to colour the vertices of G so that adjacent vertices are given different colours; that is, it's the least k so that there exists a k -colouring of \mathbf{G} . \diamond

The most basic problem you will have to complete about these is the following: given a graph \mathbf{G} , determine its chromatic number $\chi(\mathbf{G})$. Because the chromatic number is the *least* number of colours with which it is possible to colour \mathbf{G} , showing that $\chi(\mathbf{G}) = N$ will always require two steps:

1. Show that \mathbf{G} admits a colouring with $\chi(\mathbf{G}) = N - 1$ colours
2. Show that \mathbf{G} does not admit a colouring with fewer colours.

Example 4.1.3 Complete graph. What's the chromatic number $\chi(K_n)$ of the complete graph? Since every vertex is adjacent to every other vertex, any two vertices need to have different colours, and so $\chi(K_n) \geq n$. But certainly if we colour every vertex a different colour, then two adjacent vertices have the same colour, and that's a valid colouring of K_n , so $\chi(K_n) \leq n$. So $\chi(K_n) = n$. \square

Example 4.1.4 Trees. Suppose that T_n is a tree on $n \geq 2$ vertices. Then T_n has an edge, and the two vertices on this edge must be different colours, and so $\chi(T_n) \geq 2$. On the other hand, we can colour any tree with two colours as follows: pick any vertex, and colour it blue; then pick any vertex next to it and colour it red, then we can colour the vertices next to that blue, and colour the vertices next to those red, and continuing on outwards from our starting vertex. Hence, $\chi(T_n) \leq 2$ and so $\chi(T_n) = 2$.

Another way of phrasing this is that along any path we colour the vertices alternating red-blue-red-blue-red-blue. This wouldn't work for a general graph, because there may be two paths of different lengths between a pair of vertices v and w . But in trees because there is always exactly one path between any two vertices, and so once we colour one vertex, there's a unique way to colour all the others with two colours in this way. \square

Example 4.1.5 Cyclic graphs C_n . As with trees, as long as $n \geq 2$ the graph has at least one edge, and thus has $\chi(C_n) \geq 2$. Can we colour C_n with two colours?

If we could, the vertices would have to alternate red-blue-red-blue all the way around. This works if n is even, but if n is odd then the vertex we started with would have the same colour as the vertex we ended with, but they're adjacent. Thus, when n is odd we need at least three colours to colour the graph, but it's easy to do with three colours -- we can alternate red-blue-red-blue, but make the very last vertex green, for instance.

Summarizing, we have:

$$\chi(C_n) = \begin{cases} 2 & n \text{ odd} \\ 3 & n \text{ even} \end{cases}$$

□

In fact, as thinking about the examples of trees and cycles should show, we've already met one instance of the chromatic number -- a graph \mathbf{G} is bipartite if and only if $\chi(\mathbf{G}) = 2$, as follows immediately from the definition.

Another useful observation about the examples we've seen is that, since the chromatic number $\chi(\mathbf{G})$ is the *least* number of colours needed to colour $\chi(\mathbf{G})$, to show that $\chi(\mathbf{G}) = s$ requires doing two things:

1. Showing that \mathbf{G} *can* be coloured with s colours, and hence $\mathbf{G} \leq s$
2. Showing that \mathbf{G} *can't* be coloured with $s - 1$ colours, and hence $\mathbf{G} \geq s$

Example 4.1.6 The Wheel graph. The wheel graph W_n consists of an n -cycle together with one additional vertex, that is connected to all vertices of the n -cycle. Note that this with this convention, W_n confusingly has $n - 1$ vertices; other people may use a different convention where W_n has n vertices, but then it only has $n - 1$ vertices on the actual wheel.

Since the central vertex is connected to all other vertices, once we colour it, we can never use that colour again. But deleting that vertex we just have the n -cycle, and we already know the chromatic number of that. So we have $\chi(W_n) = \chi(C_n) + 1$. □

Definition 4.1.7 We write $\Delta(\mathbf{G})$ for the maximum degree of any vertex in \mathbf{G} :

$$\Delta(\mathbf{G}) = \max_{v \in \mathbf{G}} d(v)$$

◇

Theorem 4.1.8 We have $\chi(\mathbf{G}) \leq \Delta(\mathbf{G}) + 1$

Proof. We need to show that we can colour any graph \mathbf{G} with $\Delta(\mathbf{G}) + 1$ colours. But we can just colour the vertices of \mathbf{G} one by one in whatever order we want. When we go to colour the i th vertex v_i , we look at the $d(v_i)$ vertices adjacent to v_i . Some of them may not be coloured yet, in which case they don't affect anything, but for each vertex adjacent to v_i that is coloured, we can't use that colour for v_i .

So there are most $d(v_i) \leq \Delta(\mathbf{G})$ colours we have to avoid; if we have $\Delta(\mathbf{G}) + 1$ colours to choose from we can always find one that hasn't been used at a vertex adjacent to v_i . ■

Chromatic index and applications

It is a natural twist of the definition of chromatic number to try to colour the edges of a graph instead; the least number of colours needed is called the *chromatic index*. After introducing this concept and giving some examples, we give some story problem type questions that boil down to finding either the chromatic number or chromatic index.

Definition 4.2.1 Chromatic index. The *chromatic index* $\chi'(\mathbf{G})$ of a graph \mathbf{G} is the least number of colours needed to colour the edges of \mathbf{G} so that any two edges that share a vertex have different colours. \diamond

Not that as with the chromatic number, since the chromatic index $\chi'(\mathbf{G})$ is the minimum number of colours such that the edges can be coloured with adjacent edges having different colours, to show $\chi'(\mathbf{G}) = N$ typically requires two steps:


1. Prove that the edges of \mathbf{G} can be coloured with N colours
2. Prove that the edges of \mathbf{G} cannot be coloured with less than N colours

Example 4.2.2 The complete graph K_4 . Let's find $\chi'(K_4)$. Picking any vertex v , there are three edges incident to v , and none of these edges can have the same colour (as they all meet at v). Hence, we have $\chi'(K_4) \geq 3$.


On the other hand, it is easy to colour the edges of K_4 with three colours, as seen below, and so $\chi'(K_4) \leq 3$, and hence $\chi'(K_4) = 3$. \square

Example 4.2.3 The complete graph K_5 . Now, let's move on to K_5 . Again, looking at any vertex we see all the edges adjacent to that vertex must be different colours, and so we have $\chi'(K_5) \geq 4$. Let's try to colour the edges of K_5 with 4 colours.

Suppose we coloured the four edges adjacent to the top vertex blue, green, red and purple, from left to right, and now look at the bottom edge. It is adjacent to edges coloured green and red, and so must be blue or purple. By symmetry, it's equivalent to colour it either colour, so let's suppose it's blue, giving us the following picture:




Now the edge on the right is adjacent to edges coloured red, blue and purple, and so must be green. But now we have a problem -- consider the edges labeled x and y in the next drawing:



Both edges share vertices with edges coloured green, blue, and purple, and hence each would need to be coloured red. But they also share a vertex with each other, and so cannot both be coloured red. So we see $\chi'(K_5) \geq 5$.

On the other hand, it is easy to colour the edges of K_5 with 5 colours: colour each edge in the outside pentagon a different colour. For each edge in the outside pentagon there will be a unique edge in the inside star that does meet that edge (the one "parallel" to it) -- draw that edge the same colour. That results in the following colouring:



□

In the examples above, we found lower bounds for $\chi'(\mathbf{G})$ by considering the degrees of vertices; this argument easily adapts in general.

Theorem 4.2.4 For any graph \mathbf{G} we have $\chi'(\mathbf{G}) \geq \Delta(\mathbf{G})$

Proof. Let $v \in \mathbf{G}$ be a vertex of maximal degree $d(v) = \Delta(\mathbf{G})$. Then none of the $\Delta(\mathbf{G})$ edges incident to v can be coloured the same colour, and so we have $\chi'(\mathbf{G}) \geq \Delta(\mathbf{G})$ ■

It turns out that this nearly determines the chromatic index $\chi'(\mathbf{G})$ -- it can be at most one more than $\Delta(\mathbf{G})$:


Theorem 4.2.5 Vizing's Theorem. *For any graph \mathbf{G} we have $\Delta(\mathbf{G}) \leq \chi'(\mathbf{G}) \leq \Delta(\mathbf{G}) + 1$*

Proof. The lower bound was just proved in the previous theorem. The other direction is more difficult. ■

We now show how determining the chromatic number and chromatic index can show up as part of story questions.

Suppose there are six friends, Alice, Bob, Charlie, Dora, Elizabeth and Frank, and there is the following graph between them:

A
X
B
X
C
X
X
D
X
X
E
X
X
X
F



Here are two word problems related to \mathbf{G} :

1. The friends want to divide into groups, but the edges indicate people who currently annoy each other. What's the least number of groups the friends can divide into groups so that no group contains two people who annoy each other?
2. The friends want to hold a snooker tournament, with everyone playing three matches; the edges indicate pairs of friends who will play against each other. If multiple matches can be played each day, but each person can only be involved in one match a day, how many days are necessary to hold the tournament?

The first case concerns the chromatic number -- each group of people will be the people who have the same colour, and we don't want vertices with an edge between them to have the same colour.

The second case concerns the chromatic index -- the edges are the games that are being played, and all edges that are the same colour will be played on the same day.

Let us quickly compute the chromatic number and chromatic index of the graph \mathbf{G} above. To compute the chromatic number, we observe that the graph

contains a triangle, and so the chromatic number is at least 3. But it is easy to colour the vertices with three colours -- for instance, colour A and D red, colour C and F blue, and colour E and B green. So $\chi(\mathbf{G}) = 3$.

To compute $\chi'(G)$, since A has degree three we have $\chi'(\mathbf{G}) \geq 3$. On the other hand, it is easy to colour the edges with three colours -- for instance, colour AB, CE and DF red, colour AE, CD and BF blue, and colour AC, BD and EF green. So $\chi'(\mathbf{G}) = 3$ as well.

Introduction to the chromatic polynomial

For the chromatic number, we were asking whether or not it was possible to colour the vertices of \mathbf{G} with a given number of colours. The chromatic polynomial extends this question, and asks the following. Suppose we have k colours. How many different ways can we colour the vertices of \mathbf{G} ? It is *not* clear that this should be a polynomial in k , and proving that it is in fact a polynomial in k is the highlight of the section.

Definition and examples

Definition 4.3.1 The chromatic polynomial $P_{\mathbf{G}}$. The chromatic polynomial $P_{\mathbf{G}}$ of a graph \mathbf{G} is the function that takes in a non-negative integer k and returns the number of ways to colour the vertices of \mathbf{G} with k colours so that adjacent vertices have different colours. \diamond

It is immediate from the definition of the chromatic polynomial that $\chi(\mathbf{G})$ is the least positive number k such that $\chi_{\mathbf{G}}(k) \neq 0$.

It is *not* immediate from the definition of the chromatic polynomial that it is, in fact, a polynomial. In fact, proving that will take a little bit of work, and we postpone this until the end. First, we look at some examples of the chromatic polynomial; in many cases it is possible to easily compute the chromatic polynomial by working "vertex by vertex".

Example 4.3.2 The empty graph E_n . Recall that the empty graph E_n has n vertices and no edges. To compute $P_{E_n}(k)$ we need to count the number of ways to colour the vertices with k colours. But since E_n has no edges, we can colour each of the n vertices any of the k colours; the choices are completely independent. So $P_{E_n}(k) = k^n$ \square

Example 4.3.3 The complete graph K_n . Let's label the vertices v_0, \dots, v_{n-1} , and colour them one by one in the given order. When we colour the first vertex v_0 , no other vertices have been coloured, and we can use whichever of the k vertices we like. However, when we go to colour v_1 we note that it is adjacent to v_0 , and so whatever colour we used for v_0 we can't for v_1 , and so we have $k - 1$ colours to choose for v_1 .

Continuing in this way, we see that since all the vertices are adjacent, they all must have different colours. So when we go to colour v_i , we have already coloured v_0, \dots, v_{i-1} with i different colours, and we can't use any of these to colour v_i , and so we have $k - i$ choices to colour v_i .


Putting it all together, we see that:

$$P_{K_n}(k) = k \cdot (k - 1) \cdot (k - 2) \cdots k - n + 1$$

\square

Example 4.3.4 Two triangles joined at a vertex.


Consider the graph \mathbf{G} consisting of two triangles joined at the right a vertex, shown at the right. We can calculate $P_{\mathbf{G}}(k)$ by working vertex by vertex: there are k ways to colour vertex 1, but then when we colour vertex 2 it can't be the same as vertex 1, and so there are $k - 1$ ways to colour it. Vertex three is adjacent to both one and two, so there are $k - 2$ ways to colour it, and vertex 4 is adjacent to one coloured vertex, vertex 3, so there are again $k - 1$ ways to colour it, and finally vertex 5 is adjacent to vertices 3 and 4 and so we have $k - 2$ ways to colour it. Putting it together, we see $P_{\mathbf{G}}(k) = k(k - 1)^2(k - 2)^2$.



□

Example 4.3.5 Two triangles joined along an edge.

Consider the graph H consisting of two triangles joined along an edge, shown at the right. We can calculate $P_H(k)$ by working vertex by vertex: there are k ways to colour vertex 1, but then when we colour vertex 2 it can't be the same as vertex 1, and so there are $k - 1$ ways to colour it. Vertex three is adjacent to both one and two, so there are $k - 2$ ways to colour it, and vertex 4 is adjacent to vertices 2 and 3, which have different colours as they are adjacent to each other, so there are $k - 2$ ways to colour it. Putting it together, we see $P_H(k) = k(k - 1)(k - 2)^2$.



□

Gluing

What can we say about the chromatic polynomial of a graph \mathbf{G} that's the disjoint union of two smaller graphs: $\mathbf{G} = \mathbf{G}_1 \sqcup \mathbf{G}_2$?

We already covered the extreme case where $\mathbf{G} = E_n$ is just a disjoint union of vertices; we could colour each vertex independently of the others, as there were no edges between them at all. A similar argument works in general to give the following.

Proposition 4.3.6 *Let $\mathbf{G} = \mathbf{G}_1 \sqcup \mathbf{G}_2$ be a disconnected graph. Then $P_{\mathbf{G}}(k) = P_{\mathbf{G}_1}(k)P_{\mathbf{G}_2}(k)$*

Proof. A colouring of \mathbf{G} with k colours gives a colouring of \mathbf{G}_1 with k colours and a colouring of \mathbf{G}_2 with k colourings. Similarly, since $\mathbf{G}_1, \mathbf{G}_2$ are disconnected, how we colour one will have no effect on what colourings are possible for the other. Hence, colouring \mathbf{G} is exactly the same as colouring \mathbf{G}_1 and \mathbf{G}_2 ■

It gets more interesting if \mathbf{G} is *almost* the disjoint union of two graphs, by which we mean that \mathbf{G} is the union of two subgraphs \mathbf{G}_1 and \mathbf{G}_2 that are nearly disjoint -- perhaps they share a single vertex, or two vertices connected by an edge. In these cases, there are nice "gluing" formulas relating the chromatic polynomials of \mathbf{G}, \mathbf{G}_1 and \mathbf{G}_2 ; as $\mathbf{G}_1 \cap \mathbf{G}_2$ grows more complicated it is still possible to say something but the answer gets more complicated and not worth it.

Proposition 4.3.7 *Suppose \mathbf{G} has two subgraphs \mathbf{G}_1 and \mathbf{G}_2 , so that their union is all of \mathbf{G} , but their intersection is a single vertex v , i.e. $\mathbf{G}_1 \cup \mathbf{G}_2 = \mathbf{G}$ and $\mathbf{G}_1 \cap \mathbf{G}_2 = \{v\}$. Then we have the following relation between their chromatic polynomials:*

$$P_{\mathbf{G}}(k) = \frac{1}{k} P_{\mathbf{G}_1}(k) P_{\mathbf{G}_2}(k)$$

Proof. As in the proof of colourings of disjoint unions, a colouring of \mathbf{G} gives a colouring of both \mathbf{G}_1 and \mathbf{G}_2 by restriction, but we don't get any two colourings: both \mathbf{G}_1 and \mathbf{G}_2 contain v , and our two colourings must both make v the same colouring.

In the other direction, if we have colourings of \mathbf{G}_1 and \mathbf{G}_2 that have the same colour at v , it is clear that we can glue them together to get a colouring of \mathbf{G} . So the question reduces to the following: given that we want vertex v to have a fixed colour, how many colourings of \mathbf{G}_2 are there that colour v this colour?

The k colours are completely interchangeable, however; we could just change the names of each one. Thus, it is clear that there are as many colourings of \mathbf{G}_2 where v is red as there are where v is blue as there are where v is any given colour. Hence, if we have k colours to use, exactly $1/k$ of all colourings of \mathbf{G}_2 will have v any given colour, namely $P_{\mathbf{G}_2}(k)/k$.

Thus, to colour \mathbf{G} with k colours, we first could first colour \mathbf{G}_1 in one of the $P_{\mathbf{G}_1}(k)$ ways. This will give us a some fixed colour of v , and we saw above that there are $P_{\mathbf{G}_2}(k)/k$ colourings of \mathbf{G}_2 where v has this colour, and so we have the result. ■

Example 4.3.8 Two triangles joined at a vertex, again. As an example, we revisit our previous example. For \mathbf{G} , two triangles joined at a vertex, we have $\mathbf{G}_1 \cong \mathbf{G}_2 \cong C_3$. Since $C_3 \cong K_3$, we know $P_{C_3}(k) = k(k-1)(k-2)$. Thus, we have:

$$P_{\mathbf{G}}(k) = \frac{P_{C_3}(k)P_{C_3}(k)}{k} = \frac{k^2(k-1)^2(k-2)^2}{k} = k(k-1)^2(k-2)^2$$

□

Proposition 4.3.9 *Suppose \mathbf{G} has two subgraphs \mathbf{G}_1 and \mathbf{G}_2 , so that their union is all of \mathbf{G} , but their intersection is a single vertex edge e connecting two vertices v and u , i.e. $\mathbf{G}_1 \cup \mathbf{G}_2 = \mathbf{G}$ and $\mathbf{G}_1 \cap \mathbf{G}_2 = \{e\}$. Then we have the following relation between their chromatic polynomials:*

$$P_{\mathbf{G}}(k) = \frac{1}{k(k-1)} P_{\mathbf{G}_1}(k) P_{\mathbf{G}_2}(k)$$

Proof. The proof is extremely similar to that of the previous proposition. A colouring of \mathbf{G} gives us colourings of \mathbf{G}_1 and \mathbf{G}_2 , but not any two colourings: they need to match at both v and at u .

Now, v could be any of the k colours, but u , being adjacent to v , can't be the same colour, and so it has $k-1$ possibilities given a choose of colour for v . Thus, there are $k(k-1)$ ways to colour v and u , and all possibilities will occur equally often within the colourings counted by $P_{\mathbf{G}_2}(k)$. Hence, given a colouring of \mathbf{G}_1 , there will be $P_{\mathbf{G}_2}(k)/[k(k-1)]$ ways to extend that colouring to one of all of \mathbf{G} , giving us the result. ■

Example 4.3.10 Two triangles joined along an edge, again. As an example, we revisit our previous example. For H , two triangles joined along an edge, we can apply the theorem with $\mathbf{G}_1 \cong \mathbf{G}_2 \cong C_3$. Thus, we have:

$$P_H(k) = \frac{P_{C_3}(k)P_{C_3}(k)}{k(k-1)} = \frac{k^2(k-1)^2(k-2)^2}{k} = k(k-1)(k-2)^2$$


□

Case by case analysis

The methods of gluing and working vertex by vertex make many chromatic polynomials easy to calculate. Other graphs, however, are not amenable to our gluing theorems, and require considering several cases when working vertex by vertex.

Example 4.3.11 The chromatic polynomial of the cycle C_4 .

Now consider the graph C_4 , shown at right, and suppose we try to count the number of colourings of it with k colours. Vertex 1 can be any of k colours, and vertex 2 has $k-1$ possibilities -- any colour except the one used for vertex 1. Moving to vertex 4, we see it is just adjacent to 1 as well, and so has $k-1$ possibilities as well.



It becomes more difficult when we try to colour vertex 3. It is adjacent to vertices 2 and 4, and so cannot be the same colour as either of these. However, vertices 2 and 4 are not adjacent, and so we don't know whether they have the same colour or not. If vertices 2 and 4 have the same colour there are $k-1$ possibilities for vertex 3, while if vertices 2 and 4 have different colours, there are only $k-2$ possibilities. Thus, we must count how many possibilities are in each of these cases.

If we want vertices 2 and 4 to have the same colour, we can first colour vertex 1 in k different ways, and then pick any of the remaining $k-1$ colours for vertices 2 and 4. Then, to complete this to a colouring of C_4 , we have to colour v_3 , which can be any of the $k-1$ colours that aren't the colour v_2 and v_4 are coloured. Thus, the case where v_2 and v_4 have the same colour has $k(k-1)^2$ possibilities.

If we want vertices 2 and 4 to have different colours, then we can first colour v_1 any of k colours, colour v_2 any of $k-1$ colours. Now, when we go to colour vertex 4 it can't be the same colour as vertex 1 since they are adjacent, and it can't be the same colour as vertex 2 by our supposition. Vertices v_1 and v_2 have different colours, and so this leaves $k-2$ possibilities for v_4 . Thus there are $k(k-1)(k-2)$ possibilities to colour vertices 1, 2 and 4 so that 2 and 4 have different colours, and then there are $k-2$ possibilities left for vertex 3, giving $k(k-1)(k-2)^2$ ways to colour C_4 so that vertices 2 and 4 have different colours.


Adding the two cases together, this gives:

$$\begin{aligned} P_{C_4}(k) &= k(k-1)^2 + k(k-1)(k-2)^2 \\ &= k(k-1)[k-1 + (k-2)^2] \\ &= k(k-1)(k^2 - 3k + 3) \end{aligned}$$

□

Example 4.3.12 The chromatic polynomial of C_5 .

A similar case-by-case argument that we made for C_4 works for C_5 , except now we have a further case to deal with. Using the vertex labellings as shown to the right, note that Vertices 3 and 4 must have different colours. We will consider three cases: Vertices 1, 3 and 4 all have different colours, Vertex 1 has the same colour as Vertex 3, and Vertex 1 has the same colour as Vertex 4.



Case 1: 1, 3 and 4 have different colours. Then there are $k(k-1)(k-2)$ ways to colour vertices 1, 3, and 4, since they must all be different colours. When we go to colour vertices 2 and 5, we see that they are each adjacent to two vertices known to have different colours, and so there are $k-2$ ways to colour each of them. Thus in total, Case 1 contains $k(k-1)(k-2)^3$ colourings.

Case 2: 1 and 3 have the same colour. Then there are k ways to choose this colour, and $k-1$ ways to choose a colour for vertex 2 (since 1 and 3 have the same colour), and $k-1$ ways to choose a colour for vertex 4. When we colour vertex 5, we know that 1 and 4 must have different colours, and so we have $k-2$ ways to colour vertex 5. Thus in total, Case 2 contains $k(k-1)^2(k-2)$ colourings.

By symmetry, we see that Case 3, where 1 and 4 have the same colour, is identical to Case 2. Thus, putting the three cases together, we see that:

$$\begin{aligned} P_{C_5}(k) &= k(k-1)(k-2)^3 + 2k(k-1)^2(k-2) \\ &= k(k-1)(k-2)[(k-2)^2 + 2(k-1)] \\ &= k(k-1)(k-2)(k^2 - 2k + 2) \end{aligned}$$

□

Chromatic Polynomial continued

It may seem plausible that, if we consider enough cases, the case-by-case method of computing the chromatic polynomial would work for any graph, no matter how complicated. Assuming this, it would follow that the chromatic polynomial $P_G(k)$ is in fact a polynomial. However, plausibility does not make a proof. In this section we complete the proof elegantly using induction, but first need to introduce the notion of deletion and contraction.

Deletion-Contraction

Definition 4.4.1 Deletion. Let G be a graph, and let $e \in G$ be an edge. Then we use $G \setminus e$ to denote the graph with the same vertex set as G , and with all the same edges, except the edge e deleted. ◇

Definition 4.4.2 Contraction. Let G be a simple graph, and let $e \in G$ be an edge between vertices v and w . Then G/e , the graph with e contracted. More precisely, G/e is the simple graph with vertices $V(G/e) = V(G) \setminus \{v, w\} \cup e$. Two non- e vertices are adjacent in G/e if and only if they were adjacent in G , and a vertex y is adjacent to e in G if and only if it was adjacent to v or w in G . ◇

Formally, Definition 4.4.2 seems formidable, but intuitively it is not as bad as the definition looks: we are changing G by making the whole edge e into

one vertex. This may create multiple edges if a vertex was adjacent to both v and w , and if so we simply remove any duplicate edges.

Lemma 4.4.3 Deletion-Contraction. *Let \mathbf{G} be a simple graph, and let $e \in \mathbf{G}$ be any edge. Then we have:*

$$P_{\mathbf{G}}(k) = P_{\mathbf{G} \setminus e}(k) - P_{\mathbf{G}/e}(k)$$

Proof. Since \mathbf{G} and $\mathbf{G} \setminus e$ have the same vertex, it is not too difficult to compare their colourings. Any valid colouring of \mathbf{G} will give a colouring of $\mathbf{G} \setminus e$, but not every colouring of $\mathbf{G} \setminus e$ arises this way -- in \mathbf{G} , colourings require that v and w , the two endpoints of e , have different colours, but there is no such requirement for $\mathbf{G} \setminus e$. So, we want to count the colourings of $\mathbf{G} \setminus e$ where v, w have the same colour. But these are just the colourings of \mathbf{G}/e : given a colouring of \mathbf{G}/e we can get a colouring of \mathbf{G} by giving most vertices the same colour, and giving both v, w whatever colour e was. By definition, we see that v and w will have the same colour in this colouring; and given any colouring of $\mathbf{G} \setminus e$ where v, w have the same colour we can get a colouring of \mathbf{G}/e by colouring e the colour that v, w were. ■

Example 4.4.4 Chromatic polynomial of C_4 , again. Let us compute $P_{C_4}(k)$ a different way as an illustration of how deletion-contraction works. No matter which edge of C_4 we pick, $C_4 \setminus e$ will be the path graph P_4 , which is a tree, and hence we know has chromatic polynomial $k(k-1)^3$. Similarly, we have that $C_4/e = C_3$, and we know $P_{C_3}(k) = k(k-1)(k-2)$. Hence we have:

$$\begin{aligned} P_{C_4}(k) &= P_{C_4 \setminus e}(k) - P_{C_4/e}(k) \\ &= P_{P_4}(k) - P_{C_3}(k) \\ &= k(k-1)^3 - k(k-1)(k-2) \\ &= k(k-1)((k-1)^2 - (k-2)) = k(k-1)(k^2 - 3k + 3) \end{aligned}$$

□

Since $\mathbf{G} \setminus e$ and \mathbf{G}/e both have fewer edges than \mathbf{G} , it follows that we can repeatedly apply Deletion-Contraction to \mathbf{G} until we have no edges left at all, and hence that Deletion-Contraction can compute the chromatic polynomial of any graph. In practice, this can be quite a complicated and formidable way to compute, but in the next section we show that this method can elegantly prove the chromatic polynomial is always a polynomial, and in some cases give nice formulas for this polynomial.

Chromatic polynomial is a polynomial

We first prove that the chromatic polynomial is actually a polynomial, but iterative use of deletion-contraction.

Theorem 4.4.5 *Let \mathbf{G} be a simple graph with n vertices and m edges. Then $P_{\mathbf{G}}(k)$ is a polynomial of degree n in k , and moreover:*

$$P_{\mathbf{G}}(k) = k^n - mk^{n-1} + \text{lower order terms}$$

Proof. The proof follows from induction on m , the number of edges, using deletion-contraction for the inductive step.

As a base case, we take $m = 0$. Then if \mathbf{G} has n vertices it must be the empty graph E_n . We have seen that $P_{E_n}(k) = k^n$, which indeed has the correct form needed for the theorem.

Now, for the inductive step, we assume that for any graph H with $\ell < m$ edges and n vertices, we have $P_H(k) = k^n - \ell k^{n-1} + \text{lower order terms}$.

Now let \mathbf{G} be any graph with m edges. We can assume $m > 0$, as the case where $m = 0$ is the base case; this means that \mathbf{G} has at least one edge e . We apply deletion-contraction to the edge e .

If we delete e , the resulting graph $\mathbf{G} \setminus e$ still has n vertices, but it now has $m - 1$ edges. Since this is less than m , we know by the inductive hypothesis that

$$P_{\mathbf{G} \setminus e}(k) = k^n - (m - 1)k^{n-1} + \text{lower order terms}$$

If we contract e , the resulting graph \mathbf{G}/e has $n - 1$ vertices, and we don't know exactly how many edges it has (contracting it may create multiple edges that need to be deleted), but it has at most $m - 1$ edges, and so by the inductive hypothesis we know that

$$P_{\mathbf{G}/e}(k) = k^{n-1} + \text{lower order terms}$$

Thus, applying deletion-contraction we have:

$$\begin{aligned} P_{\mathbf{G}}(k) &= P_{\mathbf{G} \setminus e}(k) - P_{\mathbf{G}/e}(k) \\ &= k^n - (m - 1)k^{n-1} + \text{lower order terms} - (k^{n-1} + \text{lower order terms}) \\ &= k^n - mk^{n-1} + \text{lower order terms} \end{aligned}$$

which is what we needed to show to finish the inductive step. ■

We end by showing that sometimes the inductive method of iteratively using deletion-contraction can compute explicit formulas for the chromatic polynomials of an infinite family of graphs.

Lemma 4.4.6 Chromatic polynomial of C_n . *Let C_n be the cycle graph on n vertices. Then:*

$$P_{C_n}(k) = (k - 1)^n + (-1)^n(k - 1)$$

Proof. We will induct on n . As a base case, when $n = 3$, we have:

$$\begin{aligned}(k-1)^3 + (-1)^3(k-1) &= k^3 - 3k^2 + 3k - 1 - (k-1) \\ &= k^3 - 3k^2 + 2k \\ &= k(k-1)(k-2) \\ &= P_{C_3}(k)\end{aligned}$$

For the inductive step, we assume that the proposition holds for $n-1$ and want to prove that it holds for n . We will compute $P_{C_n}(k)$ by deletion-contraction. If we delete any edge of C_n , we get the path graph P_n , and we know

If we contract any edge of C_n , we get C_{n-1} , and we know by the inductive hypothesis that

Plugging these into deletion-contraction gives:

$$\begin{aligned}P_{C_n}(k) &= P_{C_n \setminus e}(k) - P_{C_n/e}(k) \\ &= P_{P_n}(k) - P_{C_{n-1}}(k) \\ &= k(k-1)^{n-1} - ((k-1)^{n-1} + (-1)^{n-1}(k-1)) \\ &= (k-1)^{n-1}(k-1) - (-1)^{n-1}(k-1) \\ &= (k-1)^n + (-1)^n(k-1)\end{aligned}$$

as was desired. ■